



Bayesian Surface Habitability Mapping of Titan

from the Late Heavy Bombardment to the Solar Red-Giant Phase

A report submitted for the degree of
Master of Philosophy
Planetary Science and Life in the Universe

Chris Meadows, Darwin College

cm10004@cam.ac.uk



Easter Term 2026

All code, processed data products, and static habitability maps are openly archived at

<https://github.com/chrimblechrumble/pslu-p3>.

v24 — 8th May 2026

1 Introduction

1.1 Titan as an Astrobiological Target

The search for habitable environments beyond Earth has increasingly focused on icy ocean worlds in the outer solar system. Among these, Saturn’s moon Titan occupies a singular position. It is the only moon with a dense atmosphere — dominated by molecular nitrogen at 1.467 bar surface pressure and 93.65 K (Fulchignoni et al., 2005) (Appendix A) — and the only known extraterrestrial world to host stable surface liquids (Stofan et al., 2007), completing a full precipitation–evaporation–lake cycle structurally analogous to Earth’s hydrological cycle. Its atmosphere continuously produces complex organic tholins at $\sim 5 \times 10^{-14} \text{ g cm}^{-2} \text{ s}^{-1}$ (Lavvas et al., 2011), building a global organic sediment layer over ~ 4 Gyr (Lorenz et al., 2008). At depth, a subsurface water body at ~ 55 – 80 km was inferred from the tidal Love number $k_2 = 0.589 \pm 0.150$ (Iess et al., 2012). A 2025 reanalysis of the same Cassini radiometric data found interior tidal dissipation of ~ 3 – 4 TW inconsistent with a continuous global ocean, pointing instead to a thick slushy high-pressure ice layer with isolated pockets of liquid water (Petricca et al., 2025). Tidal heating within this slushy layer maintains liquid water pockets and may generate brine conduits through the ice shell, as investigated by Crick (2026). This convergence of cycling solvent, vast organic inventory, and subsurface water makes Titan a natural laboratory for prebiotic chemistry operating at planetary scale across geological time.

1.2 The Physical Environment

The atmosphere ($\sim 94\%$ N_2 , 5.65% CH_4 ; Niemann et al. 2005) sustains a methane cycle with observed equatorial rainfall (Turtle et al., 2011), fluvial channel networks (Miller et al., 2021), and polar hydrocarbon seas (Kraken, Ligeia, and Punga Mare) containing $\sim 70,000 \text{ km}^3$ of liquid methane–ethane (Malaska et al., 2025; Mastrogioseppe et al., 2019). The lake distribution is strikingly asymmetric: the north polar region hosts hundreds of filled lakes and seas covering $\sim 1.6\%$ of the global surface, while equivalent southern latitudes are largely devoid of standing liquid (Aharonson et al., 2009). This hemispheric imbalance arises from the eccentricity of Saturn’s orbit, which makes northern summers longer and milder than southern summers, driving net methane transport from south to north over Milankovitch-like cycles of $\sim 45,000$ yr (Aharonson et al., 2009). The asymmetry is directly relevant to the habitability maps: present-epoch solvent habitability is concentrated in the liquid-rich north, but the orbital forcing implies that the southern hemisphere hosted comparable lake systems in the geologically recent past and will do so again. The surface is geologically diverse: equatorial dune fields ($\sim 17\%$ by SAR area), plains ($\sim 65\%$), labyrinth terrain, impact craters, and mountains, as classified by the Lopes et al. (2019b) global geomorphological map that provides the terrain classification underpinning this thesis.

1.3 Defining Habitability

Habitability is a term used in diverse and sometimes incompatible ways across the planetary science literature (Cockell et al., 2016). In its strictest sense it denotes a binary property — whether an environment can sustain the metabolic activity of at least one known organism — but in practice astrobiologists use it as a continuous measure of how many prerequisites for life co-occur in a given location (Méndez et al., 2021). This thesis adopts the continuous usage. The posterior probability $P(H \mid \mathbf{f})$ at each pixel quantifies the degree to which the *necessary physical and chemical preconditions* for life are simultaneously present: liquid solvent, organic substrate, chemical energy, solvent cycling, atmosphere–surface exchange, topographic and geomorphological diversity, and subsurface water access. High $P(H)$ does not imply that life exists at a site, nor that abiogenesis could occur there; it implies that more of the known prerequisites are satisfied, and satisfied more strongly, than at a low- $P(H)$ site.

Abiogenesis — the emergence of life from non-living chemistry — demands conditions beyond those captured by the habitability proxy. In aqueous systems, the minimum requirements include liquid water, a source of organic monomers, free energy to drive polymerisation, and a mechanism for compartmentalisation (Ruiz-Mirazo et al., 2014). On Titan, two fundamentally different abiogenesis scenarios are conceivable: *aqueous abiogenesis* in transient impact-melt ponds or the future red-giant ocean, where liquid water contacts accumulated organics; and *non-aqueous abiogenesis* in the present-day hydrocarbon lakes, where azotosome membranes (Stevenson et al., 2015) could provide compartmentalisation and acetylenotrophy (McKay and Smith, 2005) could provide energy, but where the low temperature (94 K) severely constrains reaction kinetics. The habitability maps produced in this thesis therefore represent a *necessary but not sufficient* condition for either form of abiogenesis: they identify where the preconditions are jointly satisfied spatially and temporally, but cannot assess whether the additional requirements for the origin of life — sustained chemical disequilibrium, information-carrying polymers, or encapsulation — are met.

1.4 Chemical Habitability Scenarios

Four habitability pathways have been proposed for Titan’s surface, each with distinct implications for abiogenesis.

1. *Methanogenesis via acetylenotrophy* (McKay and Smith, 2005; Yanez et al., 2024): the hydrogenation of acetylene, $\text{C}_2\text{H}_2 + 3\text{H}_2 \longrightarrow 2\text{CH}_4$ ($\Delta G = -334\text{ kJ mol}^{-1}$), produces methane as the metabolic end product and provides energy well above the minimum for terrestrial methanogens. McKay and Smith (2005) proposed this as a basis for methanogenic life on Titan; Yanez et al. (2024) refined the thermochemistry under Titan ocean conditions, finding acetylenotrophy energy yields of 69–78 kJ mol^{−1} C. Acetylene is supplied by UV photolysis and detected at the surface by GCMS (Niemann et al., 2005). This pathway could sustain an existing non-aqueous biosphere but does not by itself provide a mechanism for abiogenesis.

2. *Azotosome membranes* (Stevenson et al., 2015): acrylonitrile (detected by ALMA; Palmer et al. 2017) spontaneously forms stable vesicle structures in liquid methane, providing a route to cellular compartmentalisation without aqueous chemistry. Unlike the fatty acid vesicles or coacervate droplets proposed for aqueous abiogenesis (Ruiz-Mirazo et al., 2014), azotosomes represent a fundamentally different compartmentalisation solution operating in a non-polar solvent. Compartmentalisation is a key requirement for abiogenesis; azotosomes could in principle fulfil this role in a non-aqueous origin-of-life scenario, though no complete pathway from vesicle to replicator has been demonstrated.
3. *Impact-melt prebiotic chemistry*: large impacts produce transient water–ammonia–organic melt ponds persisting for 10^3 – 10^4 yr, during which Strecker synthesis yields amino acids from surface aldehydes, ammonia, and HCN (Madan et al., 2026; Neish et al., 2010). Of all present-day scenarios, this provides the most direct route to aqueous abiogenesis, since liquid water, organic monomers, and energy are simultaneously available, though the transient duration ($\sim 10^4$ yr) may be too short to complete the transition from prebiotic chemistry to self-replicating systems. See (Nicolaidis, 2026) for crater physics modelling quantitative melt-volume research. See (Castlevetro, 2026) for research relating to abiogenesis within ammonia-rich brine pockets within Titans icy crust. See (Rahim, 2026) regarding the exogenic/endogenic mixing of surface organics with excavated water-ice at crater sites and its relevance to the prebiotic chemistry.
4. *Red-giant ocean*: when solar luminosity raises Titan’s surface above the water–ammonia eutectic (176 K; Grasset and Pargamin 2005), a global ocean will form over billions of years of accumulated organics for an estimated 200–500 Myr (Lorenz et al., 1997); the temporal reconstruction in this thesis estimates a longer window of ~ 0.9 Gyr (+5.1 to +6.0 Gya; Section 3.6). This is the most favourable abiogenesis scenario modelled in this thesis: warm liquid water in sustained contact with ~ 10 Gyr of accumulated organic substrate, persisting for hundreds of millions of years — conditions broadly analogous to those invoked for early Earth.

This thesis provides the first unified quantitative characterisation of all four scenarios across geological time.

1.5 Scope and Research Questions

Previous habitability assessments of Titan have been predominantly qualitative or point-specific: the Earth Similarity Index (Schulze-Makuch et al., 2011) produces a single scalar per body with no spatial resolution; Hendrix et al. (2019) reviewed icy ocean world habitability using qualitative criteria matrices; Affholder et al. (2021) applied rigorous Bayesian inference to Enceladus plume data but at a single spatial point and epoch. None integrates the full suite of Cassini data modalities into a spatially-resolved, multi-epoch probabilistic framework. Three properties in combination distinguish this work from all prior assessments: (i) spatial resolution at ~ 20 km per pixel, resolving individual lake margins and

crater structures rather than producing a whole-body scalar; (ii) the simultaneous use of all eight principal Cassini data modalities (SAR, VIMS, CIRS, altimetry, geomorphological map, channel map, tidal gravity, crater catalogue) rather than a subset; and (iii) temporal extrapolation across ~ 10 Gyr from the LHB to the red-giant ocean, enabling a unified narrative of how habitability evolves rather than a snapshot at a single epoch. A detailed methodological comparison is given in Appendix J.

This thesis develops a Bayesian Beta-conjugate posterior framework combining eight observational proxy features from the complete Cassini dataset into a habitability probability map at 4490 m/px resolution, applied across five temporal epochs from the Late Heavy Bombardment (-3.5 Gya) to the red-giant future ($+5.9$ Gya). The Beta-conjugate approach is preferred for its analytical tractability at 6.49 million grid pixels, its naturally bounded $[0, 1]$ output domain, and the explicit uncertainty quantification provided by 95% highest-density interval credible intervals at every pixel.

The specific research questions are:

- RQ1** *What is the spatial distribution of present-day surface habitability, and which environments produce the highest scores?*
- RQ2** *How did habitability differ during the Late Heavy Bombardment, when impact-melt ponds provided transient liquid water?*
- RQ3** *How does habitability evolve across the temporal sequence from the LHB (-3.5 Gya) to the red-giant ocean ($+5.9$ Gya)?*
- RQ4** *Does the derived habitability map independently validate the Dragonfly landing site at Selk crater?*

2 Methods

2.1 Overview of the Analytical Approach

The framework proceeds in four stages: (1) extraction of eight habitability-proxy features from the Cassini observational archive (Section 2.3); (2) Bayesian conjugate updating to produce a posterior habitability probability map at each epoch (Section 2.4); (3) temporal scaling of features across five anchor epochs and 72 intermediate geological epochs from -3.8 Gya to $+6.5$ Gya (Sections 2.5–2.6); and (4) ranking of named surface sites by posterior probability. The posterior is evaluated at every pixel of the canonical equirectangular grid, producing global maps at each epoch from which habitable-zone boundaries, site rankings, and temporal trajectories can be extracted. Acronyms, notation, and the full data-source catalogue are defined in Appendix F.

The analysis is applied at 4490 m/px resolution on a canonical equirectangular grid covering the full globe (1802 pixels in latitude $\times$ 3603 pixels in longitude = 6.49×10^6 pixels), using a west-positive longitude convention consistent with all Cassini raster products. Raw GeoTIFF data products are reprojected onto this grid using `rasterio.warp.reproject` with bilinear resampling; polar inset maps use an oblique stereographic projection (Snyder, 1987) (Appendix F.4). The grid resolution corresponds to ~ 20 km at the equator, comparable to the spatial scale of terrain-class boundaries in the Lopes et al. (2019b) geomorphological map and sufficient to resolve individual lake margins, crater structures, and dune-field boundaries. Figure 2.1 provides a schematic overview.




Figure 2.1: Schematic overview of the five-stage Titan habitability analysis pipeline. All eight habitability-proxy features are extracted in Stage 3 and combined in the Bayesian posterior update of Stage 4 to produce $P(H | \mathbf{f})$ maps at each of the five temporal configurations.

2.2 Observational Data Sources

Eight Cassini data products spanning SAR imaging, near-infrared spectroscopy, thermal emission, altimetry, geomorphological classification, fluvial channel mapping, gravity, and

crater cataloguing are used (Table 2.1; Appendix H provides an overview of the Cassini–Huygens and *Dragonfly* missions). The primary spatial data sources are the SAR backscatter mosaic ($\sim 67\%$ coverage at ~ 350 m/px), the VIMS+ISS 1.59/1.27 μm band-ratio mosaic ($\sim 50\%$ coverage; Seignovert et al. 2019; blended with terrain-class scores for global coverage), and the Lopes et al. (2019b) global geomorphological classification. Processing details are in Appendix F.

Dataset	Coverage	Habitability feature
Cassini SAR mosaic	$\sim 67\%$ surface	f_1 liquid HC; f_7 geodiversity
VIMS 1.59/1.27 μm	$\sim 50\%$ surface	f_2 organic abundance
CIRS thermal mosaic	Global ($\sim 10^\circ$)	f_4 methane cycle; f_5 surf–atm
GTDE/GTDR DEM	$\sim 90\%$ (interp.)	f_6 topographic complexity
Lopes (2019) geomorph. map	$\sim 90\%$ SAR	f_2 organic; f_7 geodiversity
Miller (2021) channel map	SAR coverage	f_5 surface–atmosphere
Tidal k_2 (global)	Global scalar	f_8 subsurface ocean
Hedgepeth (2020) craters	SAR coverage	f_7 geodiversity; past epoch

Table 2.1: Cassini data sources used in the habitability framework.

2.3 Feature Extraction

Eight habitability-proxy features f_1 – f_8 are extracted from the data sources above, each normalised to $[0, 1]$ via robust percentile clipping (2nd–98th percentile; Eq. 2.1). The features are designed to capture the principal habitability-relevant physical conditions on Titan’s surface: the availability of liquid solvent (f_1), the inventory of prebiotic organic substrate (f_2), the energy supply for metabolic reactions (f_3), the intensity of solvent cycling that drives chemical transport and concentration (f_4), the strength of atmosphere–surface chemical exchange (f_5), the topographic diversity that creates distinct micro-environments (f_6), the terrain-class heterogeneity that proxy for geochemical interfaces (f_7), and the proximity to the subsurface water–ammonia ocean (f_8).

Weight assignment rationale. The weights w_i (Table 2.2) encode the relative astrobiological importance of each condition for sustaining non-aqueous biochemistry. They are assigned from physical reasoning about which conditions are *necessary*, which are *enhancing*, and which are *secondary modulators*, following the tiered habitability framework of Cockell et al. (2016) and Méndez et al. (2021):

- **Tier 1 — necessary conditions** ($w = 0.20$ – 0.25 ; total 0.65): Liquid solvent (f_1), organic substrate (f_2), and chemical energy (f_3) are the three pillars of any biochemistry (Benner et al., 2004; McKay and Smith, 2005). No known or proposed metabolism can function without all three simultaneously. Liquid solvent receives the highest individual weight ($w_1 = 0.25$) because it is the single most spatially restrictive condition on Titan’s surface — present at only $\sim 2\%$ of the globe — and because theoretical work on non-aqueous biochemistry consistently identifies solvent availability as the rate-limiting factor (Rahm et al., 2016; Stevenson et al., 2015). Organic abundance

and chemical energy are assigned equal weights ($w_2 = w_3 = 0.20$) because both are globally available but vary substantially in concentration; neither is more limiting than the other *a priori*.

- **Tier 2 — enhancing conditions** ($w = 0.08\text{--}0.15$; total 0.23): The methane cycle (f_4) and surface–atmosphere interaction (f_5) do not provide habitability prerequisites directly but strongly modulate whether the Tier 1 conditions co-occur productively. The methane cycle receives $w_4 = 0.15$ because it is the dominant agent of chemical transport, concentrating organics at lake margins and driving wet–dry cycling that promotes polymerisation (Mitchell and Lora, 2016). Surface–atmosphere interaction ($w_5 = 0.08$) captures the physical delivery mechanism — slope-driven transport, channel termini, and lake-margin evaporation — which is important but subordinate to the cycle itself.
- **Tier 3 — secondary modulators** ($w = 0.02\text{--}0.06$; total 0.12): Topographic complexity (f_6), geomorphological diversity (f_7), and subsurface ocean proximity (f_8) modulate habitability potential at fine scales but are not prerequisites. Their low weights ensure that they contribute to site differentiation within a region (e.g. distinguishing Selk crater from adjacent plains) without dominating the global ranking. Subsurface ocean proximity receives the lowest weight ($w_8 = 0.02$) because the present-day surface expression of ocean–surface coupling is highly uncertain (Hemingway et al., 2013).

Prior mean rationale. The prior mean μ_i for each feature encodes the expected global-average value of that normalised feature before examining the spatial data. High μ_i (≥ 0.4) indicates a feature that is globally abundant or active across most of Titan’s surface: organic abundance ($\mu_2 = 0.70$) because $\sim 82\%$ of the surface is organic-rich terrain (Malaska et al., 2025), and methane cycle intensity ($\mu_4 = 0.40$) because the cycle operates at all latitudes above $\sim 30^\circ$. Low μ_i (≤ 0.05) indicates a feature that is rare or spatially restricted: liquid hydrocarbon ($\mu_1 = 0.02$) because lakes cover $\sim 2\%$ of the globe, and subsurface ocean proximity ($\mu_8 = 0.03$) because direct ocean–surface conduits are unconfirmed at most locations. Intermediate μ_i ($0.25\text{--}0.35$) indicates features with moderate global presence but substantial spatial variation: chemical energy ($\mu_3 = 0.35$), surface–atmosphere interaction ($\mu_5 = 0.35$), topographic complexity ($\mu_6 = 0.25$), and geomorphological diversity ($\mu_7 = 0.30$). The global prior mean $\mu_0 = \sum_i w_i \mu_i = 0.331$ is dominated by the organic-abundance contribution ($w_2 \mu_2 = 0.14$, 42% of μ_0), reflecting the physical reality that Titan’s surface is globally organic-rich but locally solvent-poor. The weights are independently assessed via data-derived feature importances (Section 3.4.3).

$$\hat{f}_i = \text{clip}\left(\frac{f_i - f_{i,\text{p2}}}{f_{i,\text{p98}} - f_{i,\text{p2}}}, 0, 1\right) \quad (2.1)$$

f_i	Name	w_i	μ_i	$w_i\mu_i$	Scientific rationale
<i>Tier 1: necessary conditions</i>					
f_1	Liquid HC availability	0.25	0.02	0.005	Solvent prerequisite; most spatially restrictive ($\sim 2\%$ coverage; Hayes et al. 2008)
f_2	Organic abundance	0.20	0.70	0.140	Substrate prerequisite; globally abundant (82% organic terrain; Malaska et al. 2025)
f_3	Chemical energy (C_2H_2)	0.20	0.35	0.070	Energy prerequisite; indirect SAR proxy (μ tempered; McKay and Smith 2005 ; Yanez et al. 2024)
<i>Tier 2: enhancing conditions</i>					
f_4	Methane cycle intensity	0.15	0.40	0.060	Transport agent; broad mid-high-lat activity (Mitchell and Lora 2016 ; Turtle et al. 2011)
f_5	Surface-atm. interaction	0.08	0.35	0.028	Delivery mechanism; subordinate to cycle (Hayes 2016 ; Miller et al. 2021)
<i>Tier 3: secondary modulators</i>					
f_6	Topographic complexity	0.06	0.25	0.015	Micro-environment diversity; low global roughness (Corlies et al. 2017 ; Lopes et al. 2019b)
f_7	Geomorphological diversity	0.04	0.30	0.012	Geochemical interfaces; moderate global heterogeneity (Lopes et al. 2019b ; Shock and Holland 2007)
f_8	Subsurface ocean proximity	0.02	0.03	0.001	Water conduit; uncertain surface expression (Hemingway et al. 2013 ; Iess et al. 2012)
		<i>Sum</i>	0.331 = μ_0		

Table 2.2: The eight habitability-proxy features. Weight w_i , prior mean μ_i , and contribution $w_i\mu_i$ to the global prior mean are the present-epoch values used in the Bayesian update (Eq. 2.11). The sum $\mu_0 = \sum_i w_i\mu_i = 0.331$ is directly verifiable from the $w_i\mu_i$ column. Features are grouped by tier (see text).

Figure G.1 (Appendix G) shows all eight feature maps at the present epoch, illustrating their distinct spatial patterns: the polar concentration of liquid hydrocarbon (f_1), the near-global organic abundance (f_2), the topographically modulated chemical energy proxy (f_3), and the latitude-dependent methane cycle signal (f_4). Figure G.2 shows how the relative activity of each feature is modulated across geological time. Implementation details, backfill strategies, and equations for composite features are given in Appendix G; the scientific basis for each feature is summarised below.

2.3.1 Feature Descriptions

f_1 : Liquid hydrocarbon availability ($w_1 = 0.25$, $\mu_1 = 0.02$). Liquid methane–ethane is the proposed non-aqueous biochemical solvent for any surface life on Titan ([McKay, 2016](#); [McKay and Smith, 2005](#)). At 94 K, water is structurally rigid and cannot support the molecular diffusion processes fundamental to metabolism; liquid methane and ethane are stable under Titan’s surface conditions, as confirmed by the permanent north polar seas.

The feature is derived from the Cassini SAR lake mask.

Weight justification: $w_1 = 0.25$ (highest) because solvent is the single most spatially restrictive prerequisite — without liquid, the other seven features are astrobiologically inert. Theoretical models of non-aqueous biochemistry consistently identify solvent availability as rate-limiting (Rahm et al., 2016; Stevenson et al., 2015).

Prior mean justification: $\mu_1 = 0.02$ because confirmed lakes and seas cover only $\sim 1.5\text{--}2\%$ of the global surface (Hayes et al., 2008); this low prior ensures that the feature’s high weight is concentrated at the polar lake margins rather than inflating scores globally.

f_2 : Organic abundance ($w_2 = 0.20$, $\mu_2 = 0.70$). Tholins produced by UV photolysis of the N_2/CH_4 atmosphere are the dominant surface material across most of Titan and the primary chemical substrate for any habitability scenario (Meyer-Dombard et al., 2025). Laboratory hydrolysis of tholins in liquid water produces amino acids and nucleobase analogues at $\sim 0.1\%$ yield (Neish et al., 2018), confirming their prebiotic potential. The feature is derived from the VIMS $1.59/1.27\text{ }\mu\text{m}$ band ratio where coverage exists ($\sim 50\%$ of the globe), with a global offset applied to align VIMS spectral ratios with the Lopes et al. (2019b) geomorphological terrain-class scores in the overlap region; in the remaining hemisphere, terrain-class scores are used directly (Appendix D; Section G.2). This blended mode preserves per-pixel spectral variation at polar lake sites while maintaining seamless global coverage.

Weight justification: $w_2 = 0.20$ (equal to w_3) because organic substrate is a necessary condition for any biochemistry but is not more limiting than energy *a priori* — both must be present for metabolism. It receives less weight than liquid ($w_1 = 0.25$) because organics are globally abundant whereas liquid is spatially restricted.

Prior mean justification: $\mu_2 = 0.70$ (highest of all features) because $\sim 82\%$ of the surface falls into organic-rich terrain classes (dunes 17% + plains 65% ; Malaska et al. 2025), making organics the most globally pervasive habitability-relevant resource.

f_3 : Chemical energy — acetylene ($w_3 = 0.20$, $\mu_3 = 0.35$). Life requires sustained chemical free energy. The primary candidate on Titan is the hydrogenation of acetylene ($\Delta G = -334\text{ kJ mol}^{-1}$; McKay and Smith 2005), with updated energy yields of $69\text{--}78\text{ kJ mol}^{-1}\text{ C}$ under Titan ocean conditions (Yanez et al., 2024). Since no Cassini instrument directly maps surface acetylene concentration, the feature is approximated from SAR backscatter (dark surfaces indicate organic-rich terrain where acetylene accumulates; 70% weight) and topographic inversion (atmospheric condensates settle in topographic lows; 30% weight).

Weight justification: $w_3 = 0.20$ (equal to w_2) because energy is a necessary condition for any metabolism, co-equal with organic substrate. It is not weighted above f_2 despite the thermodynamic imperative because the proxy is indirect (SAR backscatter is only a correlate of acetylene concentration, not a direct measurement).

Prior mean justification: $\mu_3 = 0.35$ (moderate) because atmospheric photolysis continuously produces acetylene globally (Strobel, 2010), ensuring some energy availability everywhere, but the SAR-based proxy cannot distinguish genuine acetylene accumulation

from other causes of low backscatter, tempering the prior below the ~ 0.5 – 0.6 value that the atmospheric production rate alone would suggest.

f_4 : Methane cycle intensity ($w_4 = 0.15$, $\mu_4 = 0.40$). The methane cycle is the dominant agent of chemical transport on Titan, delivering organic material to polar lake margins, maintaining chemical gradients at shorelines, and providing seasonal wet–dry cycling that may concentrate organic reactants (Mitchell and Lora, 2016; Turtle et al., 2011). The feature is a weighted composite of three independently computed components: the meridional temperature gradient from CIRS (Jennings et al., 2019) (50%; steep gradients drive evaporation differentials), the Miller et al. (2021) global channel density map (35%; channels mark active transport corridors), and a Gaussian latitude prior centred at $\pm 45^\circ$ (15%; polar regions where the methane cycle is most active).

Weight justification: $w_4 = 0.15$ (Tier 2, highest enhancing condition) because the methane cycle is not itself a habitability prerequisite — organisms do not metabolise rainfall — but it is the primary mechanism that brings Tier 1 resources (organics, energy substrates) into contact with the liquid solvent at lake margins. On Earth, the water cycle plays an analogous role: essential for sustaining habitable conditions but distinct from the chemical prerequisites themselves. The weight is set below the Tier 1 features but above the other Tier 2/3 features to reflect this intermediate role.

Prior mean justification: $\mu_4 = 0.40$ because the methane cycle is active across a broad latitudinal range ($\gtrsim 30^\circ$ in both hemispheres; Mitchell and Lora 2016), with confirmed precipitation events at tropical latitudes (Turtle et al., 2011) indicating that even equatorial sites experience some methane-cycle activity.

f_5 : Surface–atmosphere interaction ($w_5 = 0.08$, $\mu_5 = 0.35$). The physical interface between surface and atmosphere is where gas-phase photochemical products encounter liquid solvents and organic substrates. Lake margins and fluvial channel termini are the primary loci of this exchange (Hayes, 2016; Miller et al., 2021). The feature combines topographic slope (30%; steep slopes drive physical transport), a lake-margin proximity mask (40%; evaporation zones), and the Miller et al. (2021) global channel density map (30%). The lake-margin component receives the highest within-feature weight because shorelines are where atmospheric products, surface liquids, and organic substrates are simultaneously accessible; the remaining weight is split equally between the two transport-mechanism proxies. The posterior is insensitive to this internal split given the 8% overall feature weight.

Weight justification: $w_5 = 0.08$ (Tier 2, lower) because this feature captures the *mechanism* of organic delivery rather than its chemical significance — it is subordinate to the methane cycle ($w_4 = 0.15$) that drives the delivery process. The distinction is analogous to that between precipitation (the driver, f_4) and slope runoff (the delivery mechanism, f_5) in terrestrial hydrology.

Prior mean justification: $\mu_5 = 0.35$ (moderate) because some degree of atmosphere–surface exchange occurs everywhere on Titan through direct condensation and aeolian transport, but the most intense exchange is concentrated at lake margins and channel termini — spatially restricted features that depress the global mean below the values seen at those

specific locations.

f_6 : Topographic complexity ($w_6 = 0.06$, $\mu_6 = 0.25$). Local topographic roughness at ~ 20 km scale correlates with dissected terrain that most frequently exposes water-ice-bearing crustal material (Lopes et al., 2019b). On Earth, topographic heterogeneity creates micro-environmental diversity that underpins elevated species richness at landscape boundaries; on Titan, the analogous principle is that varied topography produces diverse chemical micro-environments. The feature is the local standard deviation of GTDE elevation in a 5×5 pixel (~ 22 km) sliding window.

Weight justification: $w_6 = 0.06$ (Tier 3) because topographic complexity is a modulating factor rather than a prerequisite — a flat lake shore can be highly habitable while a rough mountain summit is not. The feature contributes to site differentiation (e.g. distinguishing dissected terrain from smooth plains within the same terrain class) but should not dominate the ranking.

Prior mean justification: $\mu_6 = 0.25$ (low-moderate) because Titan’s surface is globally smoother than Earth’s, Mars’s, or the Moon’s: the GTDE shows median relief of ~ 200 m over 20 km baselines (Corlies et al., 2017), and large portions of the dune belt and plains show near-zero local roughness in the DEM. The interpolated nature of the GTDE in data-sparse regions (where local standard deviation approaches zero) further depresses the effective global mean.

f_7 : Geomorphological diversity ($w_7 = 0.04$, $\mu_7 = 0.30$). Astrobiologically relevant chemistry is disproportionately likely at interfaces between geochemically distinct terrain types, where redox gradients, differences in organic content, and water-ice availability drive reactions that would not proceed in uniform settings (Shock and Holland, 2007). On Titan, the most important geochemical interfaces are the dune-plains boundary (organic-rich sand meets active transport channels), crater ejecta contacts (water-ice meets organics), and labyrinth-plains transitions. The feature is the normalised Shannon entropy (Shannon, 1948) of the terrain class distribution in a sliding window of radius ~ 45 km, computed from the Lopes et al. (2019b) geomorphological map.

Weight justification: $w_7 = 0.04$ (Tier 3, second lowest) because terrain-type heterogeneity is a fine-scale modulator that distinguishes sites within a region (e.g. Selk crater from adjacent Shangri-La plains) but does not constitute a habitability prerequisite. The Shock and Holland (2007) thermodynamic framework from which this feature derives was developed for terrestrial hydrothermal systems; its applicability to Titan’s cryogenic surface chemistry is plausible but unconfirmed, justifying the conservative weight. Despite the low weight, this feature is the primary discriminator at Selk crater, where four terrain classes within 45 km produce $f_7 = 0.659$ ($7.2\times$ the global median).

Prior mean justification: $\mu_7 = 0.30$ (moderate) because the six Lopes terrain classes are distributed heterogeneously: $\sim 47\%$ of pixels have ≥ 2 classes within the sliding window, giving moderate global diversity. The research in (Rahim, 2026) provides details on observational evidence and empirically supports this feature.

f_8 : Subsurface ocean proximity ($w_8 = 0.02$, $\mu_8 = 0.03$). Titan’s interior was long thought to host a global subsurface water–ammonia ocean inferred from the tidal Love number k_2 (Iess et al., 2012), though a 2025 reanalysis of Cassini radiometric data points instead to a thick slushy ice layer with isolated liquid water pockets (Petricca et al., 2025). Isolated pockets of warm water remain potentially habitable (Affholder et al., 2025), and SAR-bright annular features around impact craters provide indirect evidence of past ocean–surface contact (Wood et al., 2010). The feature combines a global floor value of 0.03 (reflecting the subsurface water signal) with localised Gaussian enhancements at confirmed annular structures.

Weight justification: $w_8 = 0.02$ (lowest) because the present-day surface expression of ocean–surface coupling is the most uncertain of all eight features. The tidal $k_2 = 0.589 \pm 0.150$ (Iess et al., 2012) originally supported a global ocean, but reanalysis now favours a slushy interior with isolated water pockets (Petricca et al., 2025); the nature and extent of subsurface liquid therefore remains uncertain. Hemingway et al. (2013) showed that the ice shell thickness (~ 50 – 100 km) makes sustained present-day cryovolcanism energetically difficult, and Neish et al. (2024b) constrained the flux of surface organics reaching any subsurface water body via impact cratering to an upper limit of $\sim 7,500 \text{ kg yr}^{-1}$ of glycine — far too low to sustain a substantial biosphere (Affholder et al., 2025). These uncertainties are reflected in the low weight $w_8 = 0.02$.

Prior mean justification: $\mu_8 = 0.03$ (lowest of all features) because the global floor represents the uniform tidal forcing, while localised enhancements are restricted to the ~ 12 confirmed SAR-bright annular features (Wood et al., 2010), affecting $< 0.1\%$ of the surface. Tidal dissipation from Saturn’s gravitational forcing heats the interior and could generate brine conduits through the ice shell whether the interior is a global ocean or a slushy layer (Petricca et al., 2025) — providing the physical basis for why this feature matters at specific locations rather than being a uniform global constant. It explains how water-bearing material reaches the surface at cryovolcanic sites like Hotei and Sotra, and why the Gaussian enhancement at those sites is physically motivated rather than ad hoc. See also Crick (2026) for modelling of tidal conduit formation.

2.4 Bayesian Inference Model

The central methodological challenge is that no single Cassini instrument directly measures habitability. Each provides a partial and indirect proxy: radar backscatter reveals surface composition and liquid surfaces; near-infrared spectroscopy constrains organic content; thermal emission maps the temperature field driving the methane cycle; altimetry constrains topography and hence organic accumulation patterns; and the geomorphological classification synthesises the collective information into terrain types whose habitability implications can be grounded in laboratory and theoretical work. As defined in Section 1.3, the model quantifies the co-occurrence of necessary preconditions for life — not the probability of life itself or of abiogenesis.

The approach taken here is to combine these partial proxies through a Bayesian inference model that produces, at each grid cell, a posterior probability of habitability $P(H \mid \mathbf{f})$ given

the vector of observed features $\mathbf{f} = (f_1, \dots, f_8)$. Habitability H is treated as a latent variable — a continuous probability in $[0, 1]$ inferred from the observational proxies.

Justification for the Bayesian approach. The use of Bayesian probability to assess habitability has strong precedent in the planetary science literature. Catling et al. (2018) established a general Bayesian framework for exoplanet biosignature assessment, advocating that the probability of life should be expressed as a Bayesian posterior combining prior knowledge with observational likelihoods; their framework maps posterior probabilities to confidence levels from “very unlikely” ($\leq 10\%$) to “very likely” (90–100%). Affholder et al. (2021) applied Bayesian model comparison directly to habitability in the Saturn system, quantifying the probability of methanogenesis on Enceladus by integrating geochemical models within a Bayesian statistical framework. Walker et al. (2018) further developed this programme by demonstrating how Bayesian priors and likelihoods could be constrained from geochemical, climatic, and evolutionary knowledge, and how repeated observations could update the posterior probability of life. Shock and Holland (2007) and Méndez et al. (2021) established that habitability should be treated as a continuous, quantitative measure rather than a binary classification — precisely the formulation that Bayesian inference is designed to handle.

Choice of the Beta-conjugate model. The specific model used here — a Beta-conjugate (Beta-Binomial) prior-posterior update — has not previously been applied to spatially resolved habitability mapping. Its selection is motivated by three properties. First, the Beta distribution is defined on $[0, 1]$ and is the canonical prior for bounded probabilities (Gelman et al., 2013, ch. 2); since $P(H)$ is by definition a probability, no alternative distributional family is more natural. Second, conjugacy between the Beta prior and the Binomial likelihood ensures that the posterior is also a Beta distribution, obtainable by simple addition of pseudo-counts to the shape parameters; this makes the update analytically tractable at 6.49 million pixels without numerical integration or Markov Chain Monte Carlo (MCMC) sampling. Third, the two-parameter Beta family is flexible enough to represent prior beliefs ranging from near-uniform (weak prior, $\kappa \rightarrow 2$) to sharply peaked (strong prior, $\kappa \rightarrow \infty$), allowing the analyst to encode the degree of prior confidence through a single, interpretable parameter.

How features enter the model. Each feature $f_i \in [0, 1]$ is treated as the fractional outcome of λw_i virtual Bernoulli trials (Appendix C.3): high feature values contribute positive pseudo-observations shifting the posterior towards higher habitability; low values do the opposite. The prior contributes κ virtual observations and the data from all eight features contribute a further λ observations. High- f sites (polar lake shores, where $f_1=1.0$) shift the posterior rightward; low- f sites (mountain terrain, where $f_1 \approx 0$ and $f_2 \approx 0.25$) shift it leftward. The resulting posterior at each pixel is a Beta distribution whose mean is a convex combination of the prior mean and the weighted feature sum, with closed-form credible intervals that make uncertainty explicit at every location. Background on the latent variable concept, conjugacy, and the Beta-Binomial model is given in Appendix C.

2.4.1 Prior Distribution

The prior distribution over habitability at any pixel is (where $\sim$ denotes "is distributed as"):

$$H \sim \text{Beta}(\alpha_0, \beta_0), \quad \alpha_0 = \mu_0 \kappa, \quad \beta_0 = (1 - \mu_0) \kappa \quad (2.2)$$

where the global prior mean $\mu_0 = \sum_i w_i \mu_i = 0.331$ (Appendix B) and the concentration parameter κ controls how strongly the prior resists updating.

Concretely, at any pixel (x, y) on the grid, before observing any Cassini features, the probability density over the habitability value $h \in [0, 1]$ is given by the probability density function of the Beta distribution, which is defined in terms of the gamma function as:

$$p_0(h) = \frac{\Gamma(\alpha_0 + \beta_0)}{\Gamma(\alpha_0) \Gamma(\beta_0)} h^{\alpha_0-1} (1-h)^{\beta_0-1} \quad (2.3)$$

where $\Gamma(\cdot)$ is the gamma function, defined for all positive real numbers by the integral

$$\Gamma(x) = \int_0^\infty t^{x-1} e^{-t} dt \quad (2.4)$$

For positive integers, this reduces to the factorial: $\Gamma(n) = (n-1)!$. For non-integer arguments such as $\alpha_0 = 1.655$, the integral must be evaluated numerically (or looked up from standard mathematical tables); in practice it is computed by the `scipy.special.gamma` library function. The key property is the recurrence relation $\Gamma(x+1) = x \Gamma(x)$, which is the continuous generalisation of $n! = n \times (n-1)!$. This is the definition of $\text{Beta}(\alpha_0, \beta_0)$ from Eq. 2.2: the ratio of gamma functions is the Beta function $B(\alpha_0, \beta_0) = \Gamma(\alpha_0) \Gamma(\beta_0) / \Gamma(\alpha_0 + \beta_0)$, and the PDF is $h^{\alpha_0-1} (1-h)^{\beta_0-1}$ divided by this normalising constant. Substituting the numerical values:

$$p_0(h) = \frac{\Gamma(5)}{\Gamma(1.655) \Gamma(3.345)} h^{0.655} (1-h)^{2.345} \approx 9.47 h^{0.655} (1-h)^{2.345} \quad (2.5)$$

The mean of this distribution is $E[h] = \alpha_0 / (\alpha_0 + \beta_0) = \mu_0 = 0.331$, and the mode (peak) is at $(\alpha_0 - 1) / (\alpha_0 + \beta_0 - 2) = 0.655 / 3 \approx 0.218$. This prior is identical at every pixel — no spatial information has yet been incorporated.

The parameterisation ensures $\alpha_0 + \beta_0 = \mu_0 \kappa + (1 - \mu_0) \kappa = \kappa$, so κ is the total number of “virtual observations” encoded by the prior — the effective sample size of prior belief. With $\kappa = 5$: $\alpha_0 = 1.655$, $\beta_0 = 3.345$. Both exceed 1, which is the condition for the Beta distribution to be *unimodal* (possessing a single peak rather than the U-shaped or J-shaped forms that arise when either parameter falls below 1). The prior standard deviation is the square root of the Beta variance $\text{Var}(X) = \alpha\beta / [(\alpha + \beta)^2(\alpha + \beta + 1)]$ (Gelman et al., 2013, p. 581, Table A.1):

$$\sigma_0 = \sqrt{\frac{\alpha_0 \beta_0}{(\alpha_0 + \beta_0)^2 (\alpha_0 + \beta_0 + 1)}} = \sqrt{\frac{1.655 \times 3.345}{25 \times 6}} \approx 0.19 \quad (2.6)$$

This wide prior ($\sigma_0 \approx 0.19$ on a $[0, 1]$ scale) reflects the deliberate choice of a small κ : five

virtual observations constitute a weak prior, meaning that at well-observed pixels the Cassini data will dominate the posterior. The value $\kappa = 5$ was selected as the smallest integer that keeps the prior unimodal ($\alpha_0 > 1$ requires $\kappa > 1/\mu_0 \approx 3.0$) whilst remaining weak enough for data dominance. The moderate prior mean $\mu_0 = 0.331$ encodes the observation that Titan's surface is globally organic-rich ($w_2\mu_2 = 0.14$, the largest single contribution) whilst liquid solvent covers only $\sim 2\%$ of the surface ($\mu_1 = 0.02$).

The likelihood sharpness $\lambda = 6$ sets the total weight of data relative to prior. Each feature f_i contributes λw_i virtual observations to the posterior update, for a total data contribution of $\lambda \sum_i w_i = \lambda = 6$ (since $\sum_i w_i = 1$). The prior/data balance is therefore $\kappa/(\kappa + \lambda) = 5/11 \approx 0.45$, placing them on approximately equal footing; $\lambda = 6$ was chosen as the nearest integer above κ that gives slight data dominance. Robustness to these choices is demonstrated in Section 4.5.




Figure 2.2: Beta-conjugate Bayesian update for three representative Titan surface environments. *Left:* the prior distribution $H \sim \text{Beta}(1.655, 3.345)$ with mean $\mu_0 = 0.331$. Shaded bands mark the hard bounds $[P_{\min}, P_{\max}] = [0.150, 0.696]$. *Right:* posterior distributions after incorporating the Cassini feature vector $\mathbf{f}$ for each site (parameters $\kappa = 5$, $\lambda = 6$; present epoch). The update sharpens and shifts each distribution in proportion to how far the site's weighted feature sum $\sum_i w_i f_i$ departs from the prior mean. Polar lake shores (cyan) shift markedly right; the equatorial crater (red) shifts left; the dune sea (amber) remains near the prior. 95% HDI values are given in the legend.

2.4.2 Posterior Update

Given observed features $\mathbf{f} = (f_1, \dots, f_8)$ at a pixel (x, y) , the posterior is obtained by Beta-Binomial conjugate updating. The prior contributes $\kappa = 5$ virtual observations and the data from all eight features contribute a further $\lambda \sum_i w_i = \lambda = 6$ observations, placing prior and data on approximately equal footing. The posterior shape parameters at pixel (x, y) are:

$$\alpha_{\text{post}}(x, y) = \alpha_0 + \lambda \sum_{i=1}^8 w_i f_i(x, y) \quad (2.7)$$

$$\beta_{\text{post}}(x, y) = \beta_0 + \lambda \sum_{i=1}^8 w_i (1 - f_i(x, y)) \quad (2.8)$$

where $f_i(x, y) \in [0, 1]$ is the value of feature i at that pixel. Note that $\alpha_{\text{post}} + \beta_{\text{post}} = \kappa + \lambda = 11$ at every pixel, regardless of the feature values.

The full posterior probability density over $h \in [0, 1]$ at pixel (x, y) is then:

$$p(h \mid \mathbf{f}(x, y)) = \frac{\Gamma(11)}{\Gamma(\alpha_{\text{post}}(x, y)) \Gamma(\beta_{\text{post}}(x, y))} h^{\alpha_{\text{post}}(x, y)-1} (1-h)^{\beta_{\text{post}}(x, y)-1} \quad (2.9)$$

Since $\alpha_{\text{post}} + \beta_{\text{post}} = 11$ is constant, the leading factor simplifies to $\Gamma(11) = 10! = 3,628,800$ divided by the product of two gamma functions that vary per pixel. The posterior is a more concentrated (narrower) Beta distribution than the prior, shifted left or right depending on whether the weighted feature sum $\sum_i w_i f_i(x, y)$ falls below or above the prior mean. The prior density (Eq. 2.3) and the posterior density (Eq. 2.9) have the same functional form — both are Beta distributions — differing only in their shape parameters. This is the defining property of conjugacy: the prior and posterior belong to the same distributional family, and the update reduces to adding pseudo-counts to the shape parameters.

As a worked example, at Freeman Lacus $(x, y) = (210.0^\circ\text{W}, 83.0^\circ\text{N})$, the weighted feature sum is $\sum_i w_i f_i = 0.716$. The posterior shape parameters are:

$$\begin{aligned} \alpha_{\text{post}} &= 1.655 + 6 \times 0.716 = 5.951 \\ \beta_{\text{post}} &= 3.345 + 6 \times (1 - 0.716) = 5.049 \end{aligned}$$

and the full posterior density at this pixel is:

$$p(h \mid \mathbf{f}_{\text{Freeman}}) = \frac{3,628,800}{\Gamma(5.951) \Gamma(5.049)} h^{4.951} (1-h)^{4.049} \quad (2.10)$$

with posterior mean $P(H) = 5.951/11 = 0.541$ and 95% HDI $[0.28, 0.82]$.

The posterior mean — the primary habitability score reported at each pixel — is:

$$P(H \mid \mathbf{f}(x, y)) = \frac{\alpha_{\text{post}}(x, y)}{\alpha_{\text{post}}(x, y) + \beta_{\text{post}}(x, y)} = \frac{1.655 + 6 \sum_i w_i f_i(x, y)}{11} \quad (2.11)$$

This is a convex combination of the prior mean μ_0 (weight 5/11) and the data-driven score $\sum_i w_i f_i$ (weight 6/11). A 95% highest-density interval (HDI) is computed at each pixel from the posterior Beta distribution, providing a credible interval that encompasses the central 95% of posterior probability mass in the narrowest possible range. Because the Beta distribution is fully characterised by its two shape parameters, which are simple sums of prior and data contributions, the HDI can be computed analytically without sampling. Since $\alpha_{\text{post}} + \beta_{\text{post}} = \kappa + \lambda = 11$ is constant regardless of the data (reflecting the fixed total number of virtual observations), the posterior mean satisfies hard bounds determined solely by the model parameters. When all features are zero ($\mathbf{f} = \mathbf{0}$, the minimum possible data), the posterior reverts to the prior floor; when all are unity ($\mathbf{f} = \mathbf{1}$, the maximum possible), the posterior reaches the ceiling:

$$P_{\min} = \frac{\alpha_0}{11} \approx 0.150, \quad P_{\max} = \frac{\alpha_0 + \lambda}{11} \approx 0.696 \quad (2.12)$$




Figure 2.3: Beta prior-to-posterior update for three representative sites. Each curve shows the Beta distribution after incorporating successive features, from the prior (dashed) through to the 8-feature posterior (solid). Shaded region = 95% HDI of the final posterior.

2.4.3 Missing Data Imputation

Where a feature is absent at a pixel (due to SAR coverage gaps affecting $\sim 33\%$ of the globe, or VIMS gaps affecting $\sim 50\%$), it is imputed with its prior mean μ_i . Substituting $f_i = \mu_i$ into the update equations adds exactly the amounts the prior already encodes for that feature, so a missing observation contributes no evidence — the formally correct Bayesian treatment of missing-at-random data (Little and Rubin, 2019, Ch. 6). The primary coverage limitation is in the south polar region, where SAR, VIMS, and altimetry coverage all degrade simultaneously; here the posterior reverts most strongly towards the prior.

No published method exists for interpolating SAR backscatter intensity across Titan’s unmapped regions; unlike topography, which varies smoothly and can be splined from sparse profiles (Corlies et al., 2017; Lorenz et al., 2013), backscatter depends on local roughness, composition, and viewing geometry and lacks the spatial autocorrelation needed for meaningful gap-filling. The pipeline therefore assigns physically motivated floor values in SAR-unmapped pixels rather than interpolated estimates.

2.4.4 Feature Importance Validation

To assess which features provide the strongest geographic discrimination in the habitability maps, a calibrated Gaussian Naive Bayes classifier is trained as a post-hoc diagnostic on binary labels derived from the Bayesian posterior by median split. Feature importances are computed by mean-replacement sensitivity: each feature is replaced with its global mean value and the resulting change in classifier accuracy is recorded. The prior weights w_i determine the posterior exactly through the analytical formula (Eq. 2.11); the classifier importances instead reveal which features provide the strongest *spatial discrimination* between above-median and below-median sites, which may differ from the prior weights if some features have greater geographic variance than others. Results are reported in Section 3.4.3.

2.5 Temporal Configurations

2.5.1 Design Philosophy

The habitability framework is applied at five distinct temporal configurations, each representing a qualitatively different physical state of Titan. At each anchor epoch, the eight features are given physically motivated epoch-specific scale functions (Table 2.3).

The specific epoch dates are determined by physical events:

Past (−3.5 Gya). LHB peak. No polar seas ($s_1=0.10$); organic inventory at $\sim 40\%$ of present ($s_2=0.125$) after only ~ 0.5 Gyr of UV photolysis; subsurface ocean $2.5\times$ closer to the surface ($s_8=2.5$) due to reduced ice-shell thickness (Tobie et al., 2005); acetylene energy $\sim 2\times$ present ($s_3\approx 2.1$) from the young Sun’s higher UV output (Bahcall et al., 2001). The impactor flux during the LHB was $100\text{--}1000\times$ the present background rate (O’Brien et al., 2005), producing transient water–ammonia melt ponds persisting for $10^3\text{--}10^4$ yr at crater sites — the sole source of liquid water on Titan’s surface.

Lake Formation (−1.0 Gya). Polar lake consolidation epoch. Liquid hydrocarbon at the onset of sustained pooling ($s_1=0.10$, beginning the ramp towards present-day coverage); organic accumulation at $\sim 75\%$ of present ($s_2=0.75$); methane rain cycle strengthening but not yet at full intensity ($s_4=0.80$). The impact-melt proximity feature has largely faded as the LHB ended ~ 2.5 Gyr earlier. This is the period of most rapid change in the global habitability structure: the polar lake system transitions from partial to near-present coverage.

Present (Cassini era, 2004–2017). All features at directly measured Cassini values; no temporal scaling applied. This is the observational anchor of the entire framework and the epoch with highest confidence, as the posterior probabilities are derived entirely from data rather than modelled extrapolations. The fully established polar hydrocarbon sea system, active methane cycle, and mature organic inventory produce the most spatially differentiated habitability map of any epoch.

Near Future (+0.25 to +1.2 Gya; stable plateau). Solar luminosity rises from $+1.4\%$ to $\sim +6\%$ above present across this window (Bahcall et al., 2001), corresponding to surface temperature increases of $\sim 1\text{--}5$ K — entirely absorbed by the model and producing no spatially structured change in any feature. The representative animation frame is drawn from $+1.19$ Gya within the plateau; this epoch and all intermediate epochs within the window are functionally equivalent for habitability purposes. The spatial correlation between any two epochs in this range exceeds $r = 0.999$; all rankings are preserved. This demonstrates that Titan’s lake–organic habitability system is robust to main-sequence solar evolution over at least ~ 1 Gyr.

Future (+5.9 Gya). Red-giant ocean peak at $T_s \approx 466$ K, well above the water–ammonia eutectic at all latitudes. Both f_1 (liquid availability) and f_8 (subsurface ocean proximity)

are overridden to 1.0 globally, reflecting the global ocean. Organic accumulation reaches $2.5\times$ present after ~ 10 Gyr of UV photolysis. The methane cycle (f_4) is suppressed (the hydrocarbon is replaced by a water cycle) and acetylene photochemistry (f_3) is reduced under the water–ammonia atmosphere because UV penetration is inhibited. The ocean window spans +5.1 to +6.0 Gya (~ 0.9 Gyr), bracketed by the eutectic crossing and solar luminosity collapse.

2.6 Multi-Epoch Habitability Reconstruction

2.6.1 Temporal Scaling of Features

The multi-epoch reconstruction applies epoch-specific scale functions to the present-epoch feature maps:

$$f_i(t) = \text{clip}[s_i(t) \cdot f_i(0), 0, 1] \quad (2.13)$$

where $s_i(t)$ is the scale multiplier for feature i at epoch t (Table 2.3). Two time constants appear in the scale functions: $t_{\odot} = 4.57$ Gya, the current age of the Sun (Bahcall et al., 2001), which governs UV-dependent features (s_3 , s_1 at red-giant epochs); and $t_{\text{atm}} = 4.0$ Gyr, the estimated duration of active tholin photochemistry in Titan’s atmosphere, which governs organic accumulation (s_2). The posterior is evaluated at each pixel using the scaled features across 72 epochs from -3.8 Gya to $+6.5$ Gya, with PCHIP interpolation providing smooth transitions.

Feature	Scale function $s_i(t)$	Key references
f_1 liquid HC	0.10 for $t < -1.0$ Gya (pre-lake era); ramps 0.1 $\rightarrow$ 1.0 from -1.0 to -0.5 Gya (lake establishment); 1.0 until $+4.0$ Gya; declines linearly to 0 by $+5.0$ Gya (solar evaporation); overridden to 1.0 (global ocean) at $T_s \geq 176$ K.	Hayes (2016); Lorenz et al. (1997); Mitchell and Lora (2016); Tobie et al. (2006)
f_2 organic	$\min((t_{\text{atm}}+t)/t_{\text{atm}}, 2.5)$: linear accumulation of UV-photolysed tholins proportional to atmospheric age since ~ -4.0 Gya; capped at $2.5\times$ present. <i>Declared approximation</i> : a constant production rate is assumed; the actual rate was higher under the younger, more UV-active Sun (see s_3).	Cable et al. (2012); Lavvas et al. (2011); Niemann et al. (2005); Vuitton et al. (2007)
f_3 acetylene	$(t_{\odot}/(t_{\odot}+t))^{0.5}$: UV-driven C_2H_2 production scales with solar EUV/FUV luminosity; young Sun was more active. Overridden to 0.10/0.35 under global ocean (UV shielded).	Bahcall et al. (2001); Lavvas et al. (2011); Loison et al. (2015); Strobel (2010)
f_4 methane	0.60 for $t < -1.0$ Gya (episodic outgassing, pre-lake); 0.80 for -1.0 to -0.5 Gya; 1.0 until $+4.5$ Gya; declines to 0 by $+5.0$ Gya (Jeans escape and photolytic loss). Overridden to 0.30/0.40 under global ocean.	Flasar et al. (2005); Mitchell and Lora (2016); Niemann et al. (2005); Strobel (2010); Tobie et al. (2006)
f_5 surface atmosphere	$0.40 + 0.60 s_1(t)$: fixed slope-transport component (0.40); lake-margin exchange (0.60) scales with liquid coverage.	Hayes (2016); Mitchell and Lora (2016); Radebaugh et al. (2011)
f_6 topography	1.30 for $t < -2.0$ Gya; 1.15 for -2.0 to -1.0 Gya; 1.0 thereafter (crater density reduces roughness over time).	Hedgepeth et al. (2020)
f_7 geo diversity	0.70 for $t < -3.0$ Gya; 0.85 for -3.0 to -2.0 Gya; 1.0 thereafter (terrain classes consolidate).	Lopes et al. (2019b)
f_8 ocean	2.50 for $t < -2.0$ Gya; 1.80 for -2.0 to -1.0 Gya; 1.30 for -1.0 to -0.5 Gya; 1.0 thereafter; overridden to 1.0/0.03 when $T_s \geq 176$ K.	Iess et al. (2012); Nimmo and Bills (2010); Tobie et al. (2005)

Table 2.3: Temporal scale functions $s_i(t)$ applied to each feature at epoch t (Gyr from present), with supporting references. A multiplier of 1.0 means no change from the present-epoch value; values < 1 reduce the feature, values > 1 enhance it. All outputs are clipped to $[0, 1]$ after scaling. The eutectic override ($T_s \geq 176$ K) applies to f_1 , f_3 , f_4 , and f_8 at the red-giant ocean epoch (Grasset and Pargamin, 2005; Lorenz et al., 1997; Tobie et al., 2005).

The key scale functions encode the following physical transitions:

- **Liquid hydrocarbon** (s_1): held at 0.10 for $t < -1.0$ Gya, reflecting the Tobie et al. (2006) model in which sustained lake coverage is geologically recent. The ramp from -1.0 to -0.5 Gya corresponds to consolidation of Kraken and Ligeia Mare (Hayes, 2016; Mitchell and Lora, 2016). Decline to zero by $+5.0$ Gya is driven by increasing solar luminosity vaporising the seas (Lorenz et al., 1997).
- **Organic abundance** (s_2): the linear form $\min(t_{\text{elapsed}}/t_{\text{atm}}, 2.5)$ approximates steady-state UV photolysis accumulating tholins since ~ -4 Gya (Niemann et al., 2005; Tobie et al., 2006). The constant production rate underestimates early-epoch accumulation, since UV flux was higher under the younger Sun (see s_3).

- **Chemical energy** (s_3): the $(t_{\odot}/(t_{\odot}+t))^{0.5}$ scaling follows the standard luminosity evolution of Bahcall et al. (2001), applied to UV-driven acetylene production (Loison et al., 2015; Strobel, 2010).
- **Subsurface ocean** (s_8): the $2.5\times$ enhancement at $t < -2.0$ Gya reflects the thinner ice shell during early epochs, when any subsurface water body was closer to the surface and cryovolcanic conduits were potentially more active (Iess et al., 2012; Nimmo and Bills, 2010; Tobie et al., 2005). Recent evidence that the present-day interior may be slushy rather than a continuous ocean (Petricca et al., 2025) does not invalidate s_8 at past epochs, where the ice shell was thinner and warm-water pockets were likely more extensive; it does caution against treating the feature as implying a well-connected global ocean at any epoch.

At +5.1 Gya surface temperature crosses the water–ammonia eutectic (176 K; Grasset and Pargamin 2005). Liquid hydrocarbon gives way to a global water–ammonia ocean ($f_1 \rightarrow 1.0$ everywhere); subsurface ocean merges with the surface ($f_8 \rightarrow 1.0$); organic accumulation reaches $2.5\times$ present. The methane cycle becomes a water cycle; acetylene photochemistry is suppressed (Lorenz et al., 1997).

2.6.2 Validation

The reconstruction is validated against three criteria: (i) each anchor epoch’s scaled posterior is within 0.02 of the corresponding static-pipeline posterior at the 95th percentile, confirming that the scale functions faithfully reproduce the anchor-epoch feature values; (ii) the global mean rises monotonically through the lake-formation transition (-1.0 to -0.5 Gya), consistent with the physical expectation that the establishment of the polar sea system represents a genuine increase in surface habitability potential; and (iii) the N-polar/equatorial ratio follows the expected pattern: ~ 1.45 at the LHB (moderate polar enhancement from VIMS organic abundance, no polar seas), rising to ~ 1.53 at the present epoch (organic + lake dominance), and declining to ~ 1.39 at the global ocean (uniform liquid coverage reduces but does not eliminate spatial differentiation, because the organic accumulation advantage of the polar basins persists). All three criteria are satisfied. The reconstruction also reproduces the known qualitative transitions: the impact-crater signal fading between -3.0 and -2.0 Gya, the cryovolcanic signal persisting through the mid-Titan epoch, and the abrupt onset of the global ocean at the eutectic crossing.

2.7 Scope of This Thesis

This thesis uses the Bayesian framework described above to address four research questions (Section 1.5): the spatial distribution of present-day habitability (RQ1), the role of impact structures during the LHB (RQ2), the temporal trajectory from -3.5 Gya to $+5.9$ Gya (RQ3), and the independent validation of the *Dragonfly* landing site at Selk crater (RQ4). The framework is intentionally conservative: a weak prior ($\kappa = 5$) ensures data dominance; the independence assumption is acknowledged and bounded by sensitivity analysis; and all

temporal extrapolations are explicitly flagged as model-dependent. The goal is not to claim that specific locations on Titan are inhabited, but to provide a quantitative, reproducible, and updatable habitability map that can guide future mission planning and constrain the interpretation of in-situ measurements.

3 Results

3.1 Overview and Temporal Habitability Trends

The Bayesian framework produces posterior mean habitability maps $P(H \mid \mathbf{f})$ at five temporal anchor epochs and, via the multi-epoch reconstruction (Section 2.6), at 72 intermediate epochs spanning -3.8 to $+6.5$ Gya. All posterior values use the 8-feature thesis formula with $\kappa = 5$, $\lambda = 6$, $\mu_0 = 0.331$, giving hard posterior bounds $P_{\min} = 0.150$ and $P_{\max} = 0.696$ when all features lie in $[0, 1]$; the *Future* epoch exceeds $P_{\max}$ because the $2.5\times$ organic-accumulation temporal scaling pushes the effective feature sum above unity (maximum $P(H) = 0.829$). The 95% HDI for Selk crater (the most-studied individual site) is $[0.04, 0.49]$, representative of the wide credible intervals arising from the moderate information content of the Cassini observables. The five epochs are presented chronologically below.

Figure 3.1 shows the temporal trajectory. The global median rises from 0.27 at the LHB to 0.30 at *Lake Formation* (-1.0 Gya), then increases to 0.32 at the *Present* epoch as the polar lake system fully establishes. This level is maintained through the *Near Future* ($+0.25$ Gya), before a transient minimum at $\sim +5.0$ Gya when the polar seas have evaporated but the water–ammonia eutectic has not yet been crossed — a ~ 100 Myr interval during which Titan is transiently solvent-free. The global mean then rises to its maximum of 0.53 at the *Future* epoch ($+5.9$ Gya) as the red-giant ocean covers the surface.

The north polar/equatorial ratio tracks the dominance of the liquid-hydrocarbon feature superimposed on the VIMS organic-abundance signal: ~ 1.45 at LHB (organic lead alone, no polar seas), rising to ~ 1.53 at *Present* (organic + lake dominance), and declining to ~ 1.39 at the *Future* global ocean (uniform liquid, but polar organic advantage persists). The trajectory is non-monotonic, passing through three qualitatively distinct regimes: (i) organic-accumulation dominance during the LHB, where polar proto-basins lead on VIMS-derived f_2 alone; (ii) polar-lake dominance from -1.0 Gya through the present, where the liquid-hydrocarbon feature amplifies the existing polar lead; and (iii) global ocean dominance from $+5.1$ Gya, where $f_1=1.0$ everywhere and the residual polar lead is driven by the $2.5\times$ -amplified organic accumulation.




Figure 3.1: Global and regional mean posterior habitability $P(H | \mathbf{f})$ as a function of geological time from the LHB peak (−3.8 Gya) to the post-red-giant refreezing epoch (+6.5 Gya). Black solid line: global mean. Blue dashed: north polar region ($>60^\circ\text{N}$). Orange dash-dot: equatorial belt ($|\text{lat}| < 30^\circ$). Red dotted: crater sites. Vertical dashed lines mark key epoch boundaries; blue shading indicates the red-giant ocean window (+5.1–6.0 Gya).

The maps in this chapter quantify **present-day solvent habitability** — the capacity of each location to support liquid-solvent chemistry. At the present epoch, they correctly identify north polar lake shores as the globally highest-scoring environments. The *Dragonfly* mission targets a different metric — **prebiotic chemistry habitability** (past liquid-water contact) — for which impact craters are the primary targets. The relationship between these two frameworks is discussed in Section 1.2.

3.2 Past Epoch (−3.5 Gya): Late Heavy Bombardment

At −3.5 Gya, Titan was a profoundly different world. The polar hydrocarbon sea system did not yet exist ($s_1 = 0.10$), the organic inventory had accumulated for only ~ 0.5 Gyr of UV photolysis ($s_2 = 0.125$; this scale factor, when applied to the present-day feature distribution, corresponds to $\sim 40\%$ of the present absolute organic abundance due to the non-uniform feature mapping), and the methane cycle was episodic and weak.

In compensation, three factors were substantially elevated. First, the young Sun’s higher UV output drove an acetylene-energy proxy $\sim 2\times$ the present value ($s_3 \approx 2.1$), providing more chemical energy per unit surface area. Second, subsurface ocean proximity was $2.5\times$ elevated ($s_8 = 2.5$) due to reduced ice-shell thickness, bringing any subsurface water body closer to the surface. Third, the impactor flux during the LHB was $100\text{--}1000\times$ the present rate (O’Brien et al., 2005), producing transient water–ammonia–organic melt ponds at confirmed crater sites. These ponds — $\sim 50\text{--}100\text{ km}^2$ area, persisting for $10^3\text{--}10^4$ yr — provided the sole source of liquid water on Titan’s surface, during which all precursors for Strecker amino acid synthesis were simultaneously present: liquid water–ammonia solvent, aldehydes from tholin decomposition, and HCN from the nitrogen-rich atmosphere.

The *Past* map (Figure 3.2) retains a recognisable polar enhancement even at this early epoch, albeit far weaker than at the *Present*. Without surface liquid ($s_1 = 0.10$), the

polar basins lack the liquid-hydrocarbon advantage (f_1 , $w_1 = 0.25$) that makes them rank highest at later epochs; instead, the polar lead is driven entirely by the VIMS-derived organic abundance feature (f_2), which reveals that the future lake basins already host the highest organic accumulation on Titan’s surface. This signal lifts the proto-basin sites above the equatorial dune belt ($f_2 = 0.78$ – 0.82 by terrain class, but lower by VIMS spectral measurement at the specific equatorial sites).

Impact structures (where the impact-melt proximity feature creates enhancements of 0.05 – 0.15 above regional background for typical mid-size impactors; Selk’s enhancement of ~ 0.015 is near the lower end because its equatorial organic supply is below average) and cryovolcanic candidate sites (Hotei Regio, Sotra Facula) remain astrobiologically significant but no longer lead the ranking.

The top-10 scores span $P(H) = 0.331$ – 0.464 (Table 3.1), with a spread of 0.133 . All ten sites are polar proto-basins — locations that will become lake shores at later epochs — ranked here by the VIMS-derived organic abundance (f_2) accumulated at these locations, combined with elevated subsurface ocean proximity ($s_8=2.5$). Freeman Lacus (the overall #1) has the highest VIMS organic signature of any polar site. Equatorial dune fields (Belet, Shangri-La) rank outside the top 10 because the VIMS mosaic resolves per-pixel organic accumulation that favours the polar basins. Cryovolcanic sites (Hotei, Sotra) similarly drop below rank 10. In the 8-feature formula, large craters still rank below the polar sites (Selk $P(H)\approx 0.258$) because their organic abundance ($f_2 = 0.30$ – 0.42) remains substantially lower. Despite these modest scores, the impact structures are the most astrobiologically significant sites at this epoch because they are the *sole source* of transient liquid water on Titan’s surface — a distinction that the present framework, which quantifies solvent habitability rather than prebiotic chemistry potential, does not fully capture. Selk crater’s scientific importance lies in its unique position at the Shangri-La dune–plains boundary, where the organic substrate available for Strecker synthesis in the transient melt pond is maximised (Neish et al., 2018).

If life exploited these transient oases, it would resemble the most opportunistic of terrestrial extremophiles: thermophilic chemotrophs gorging on dissolved tholins in temporary warm puddles, racing to complete a metabolic cycle before their habitat freezes shut — less a biosphere than a series of biochemical flash mobs, each lasting a few thousand years before the lights go out.




Figure 3.2: Past-epoch (−3.6 Gya; LHB) posterior mean habitability map. A moderate north polar enhancement is visible, driven by VIMS organic abundance at proto-basin sites. Same colour scale as Figure 3.4.

Rank	Location	Lon (°W)	Lat (°N)	$P(H)$	Dominant features
1	△ Freeman (proto-basin)	210.0	+83.0	0.464	High VIMS f_2 ; $s_8=2.5$
2	△ Oib (proto-basin)	220.0	+82.0	0.432	High VIMS f_2 ; $s_8=2.5$
3	△ Koitere (proto-basin)	154.0	+73.0	0.432	High VIMS f_2 ; $s_8=2.5$
4	△ Ontario (proto-basin)	179.0	−72.0	0.352	South polar; VIMS f_2 elevated
5	△ Hammar (proto-basin)	93.5	+70.7	0.349	VIMS f_2 ; $s_8=2.5$
6	△ Mackay (proto-basin)	262.8	+77.5	0.346	VIMS f_2 ; $s_8=2.5$
7	△ Paxsi (proto-basin)	243.0	+72.0	0.344	VIMS f_2 ; $s_8=2.5$
8	△ Neagh (proto-basin)	327.0	+72.0	0.337	VIMS f_2 ; $s_8=2.5$
9	△ Bolsena (proto-basin)	12.3	+75.4	0.331	VIMS f_2 ; $s_8=2.5$
10	△ Logtak (proto-basin)	197.5	+74.0	0.331	VIMS f_2 ; $s_8=2.5$
	★ Selk Crater	199.0	+7.0	0.258	High f_7 ; low f_2 limits score
	★ Huygens Site	192.3	−10.6	0.239	Equatorial plains; $f_3(t)$ boosted

Table 3.1: Top-10 habitable locations at the *Past* epoch (−3.5 Gya). All are future lake-shore sites (proto-basins); ★ = lander site (below line: outside top 10). $P(H)$ from Eq. 2.11.

3.3 Lake Formation Epoch (−1.0 Gya): Polar Lakes Establishing

At −1.0 Gya, Titan is transitioning between two habitability regimes. The polar lake system is just beginning to consolidate: liquid hydrocarbon is at only 10% of present coverage ($s_1 = 0.10$), corresponding to incipient pooling in the deepest polar basins. Organic accumulation has reached ~75% of present ($s_2 = 0.75$), and the methane cycle is becoming established but has not yet reached its present intensity ($s_4 = 0.80$). The impact-melt signal has largely faded as the LHB ended ~2.5 Gyr earlier. The transition from −1.0 to −0.5 Gya shows the highest pixel-wise posterior change rate in the entire reconstruction as the polar lakes fill, converting the post-LHB declining trend into the increasing trend that reaches the present-epoch plateau.

Comparing the *Lake Formation* map (Figure 3.3) with the *Past* epoch (Figure 3.2) reveals two changes. First, the north polar lake margin signal strengthens as incipient pooling in the deepest basins adds the liquid-hydrocarbon feature ($f_1=0.10$) to the already high VIMS-derived organic signal, reinforcing the polar lead that was present even at the *Past* epoch. Second, the impact-melt crater signal has largely faded to background.

All ten top-10 sites at this epoch are polar proto-lake-shores (Table 3.2), continuing the pattern from the *Past* epoch: VIMS-derived organic abundance dominates the ranking even before the polar seas have filled. The score spread has widened to 0.125 ($P(H) = 0.364$ – 0.489), reflecting the combined effect of f_2 variation and the emerging liquid signal. Equatorial dune fields (Belet, Shangri-La) rank outside the top 10 because the VIMS mosaic resolves lower per-pixel organic abundance at these specific equatorial locations than their terrain class implies. As the polar seas fill between −0.5 Gya and the *Present* epoch, the polar advantage amplifies further rather than inverting: the liquid-hydrocarbon feature adds to the existing VIMS organic lead, widening the polar–equatorial contrast to its maximum.

Any organism inhabiting these nascent lakes would face a solvent 180 K colder than the coldest terrestrial habitat, with reaction kinetics approximately 10^6 times slower than in water at room temperature (Schulze-Makuch and Grinspoon, 2005). Life here would not so much swim as *diffuse* — a metabolism measured in centuries per cell division, making even the most patient terrestrial glacier microbe look hyperactive by comparison.




Figure 3.3: Lake-formation epoch (−0.4 Gya) posterior mean habitability map. Polar proto-lake-shores dominate the ranking; the VIMS organic signal strengthens as incipient pooling adds f_1 . Compare with Figure 3.2 and Figure 3.4.

Rank	Location	Lon (°W)	Lat (°N)	$P(H)$	Dominant features
1	△ Freeman Lacus	210.0	+83.0	0.489	f_2 peak; $f_1=0.10$ emerging
2	△ Oib Lacus	220.0	+82.0	0.461	f_2 high; proto-lake forming
3	△ Koitere Lacus	154.0	+73.0	0.459	f_2 high; proto-lake forming
4	△ Ontario Lacus	179.0	−72.0	0.392	South polar; $f_1=0.10$
5	△ Mackay Lacus	262.8	+77.5	0.384	f_2 elevated; $f_1=0.10$
6	△ Hammar Lacus	93.5	+70.7	0.384	f_2 elevated; $f_1=0.10$
7	△ Paxsi Lacus	243.0	+72.0	0.383	f_2 elevated; $f_1=0.10$
8	△ Neagh Lacus	327.0	+72.0	0.381	f_2 elevated; $f_1=0.10$
9	△ Bolsena Lacus	12.3	+75.4	0.370	VIMS f_2 ; basin forming
10	△ Cardiel Lacus	128.0	+78.0	0.364	VIMS f_2 ; basin forming
	★ Selk Crater	199.0	+7.0	0.248	f_7 high but f_2 low
	★ Huygens Site	192.3	−10.6	0.247	Plains organic; $f_3(t)=1.3\times$

Table 3.2: Top-10 habitable locations at the *Lake Formation* epoch (−1.0 Gya). △ = lake/sea shore; ★ = lander site (below line). $P(H)$ from Eq. 2.11.

3.4 Present Epoch (Cassini Era, 2004–2017)

The *Present* epoch is the observational anchor of the entire framework: all eight features are at their directly measured Cassini values, with no temporal scaling applied. This is the only epoch at which the posterior probabilities are derived entirely from observational data rather than from modelled extrapolations, and consequently the results presented here carry the highest confidence of any epoch. Titan at the Cassini epoch possesses a fully established polar hydrocarbon sea system containing $\sim 70,000 \text{ km}^3$ of liquid hydrocarbons (Malaska et al., 2025; Mastrogioiuseppe et al., 2019), an active methane cycle with observed equatorial rainfall (Turtle et al., 2011) and polar storm systems, and an organic tholin inventory accumulated over the full ~ 4 Gyr of atmospheric photolysis history.

3.4.1 Habitability Map

The present-epoch map (Figure 3.4) is dominated by a pronounced north–south asymmetry driven by liquid hydrocarbon distribution. The highest posteriors form a crescent-shaped band at $60\text{--}85^\circ\text{N}$ along the polar lake margins, with the smaller lacus (Freeman, Oib, Koitere) achieving higher scores than the three large maria (Kraken, Ligeia, Punga) because the VIMS spectral data reveals higher per-pixel organic accumulation at the lacus sites. A secondary zone of intermediate habitability occupies the equatorial dune belt ($P(H) \approx 0.22\text{--}0.28$), where the dominant driver is organic abundance: dune sands are tholin-dominated (Le Mouélic et al., 2019), though the VIMS-blended mosaic resolves substantial variation in absolute organic values across equatorial sites. Chemical energy (f_3) is also elevated because SAR-dark backscatter is consistent with high organic content and flat interdune surfaces concentrate acetylene in topographic lows. The key limitation is near-zero surface liquid ($f_1 \approx 0.02$); episodic methane rainfall has been observed (Turtle et al., 2011) but evaporates within days to weeks. The two named equatorial sites fall squarely within this range: Selk crater ($P(H) = 0.243$) and the *Huygens* landing site ($P(H) = 0.260$) both

lie within the stated 0.22–0.28 band, confirming the representativeness of these landmarks for equatorial habitability conditions.

The lowest-scoring regions are the Xanadu bright terrain ($P(H) \approx 0.18$; water-ice-rich composition assigns the mountain terrain class $f_2 = 0.25$, a 3–4 $\times$ reduction in organics relative to adjacent dune terrain), mountain ranges (Mithrim, Doom, Erebor Montes), and the south polar region outside Ontario Lacus (suppressed methane cycle at the Cassini epoch’s northern summer). The contrast with earlier epochs is quantitative rather than qualitative under the VIMS-blended scoring: polar proto-basin sites already led at the *Past* and *Lake Formation* epochs (Tables 3.1–3.2), but the addition of the liquid-hydrocarbon feature (f_1) at the *Present* epoch widens the score spread from 0.133 (*Past*) and 0.125 (*Lake Formation*) to 0.148, and raises the absolute maximum from 0.489 to 0.541. The north polar/equatorial ratio reaches ~ 1.53 — the maximum spatial contrast of any epoch.




Figure 3.4: Present-epoch (Cassini era) posterior mean habitability map at 4490 m/px. Colour scale: 0.10 (dark; low) to 0.75 (bright; near posterior upper bound). North polar lake surfaces reach $P(H) \approx 0.42$ –0.54, shore margins ≈ 0.39 –0.54, equatorial organic terrain ≈ 0.20 –0.30. Selk crater marked (star) at 199°W, 7°N.

3.4.2 Top-10 Habitable Locations

All ten top-10 sites are lake or lacus margins (Table 3.3) — a pattern that, under the VIMS-blended organic scoring, now holds at *every* epoch rather than emerging only at the Present. The pattern arises from the coincidence of high VIMS-derived organic abundance (f_2), maximum liquid hydrocarbon ($f_1=0.72\text{--}0.85$), active methane cycle ($f_4\approx0.58\text{--}0.62$), and elevated shoreline geodiversity ($f_7\approx0.60\text{--}0.76$) that only occurs along confirmed hydrocarbon-filled coastlines.

Freeman Lacus achieves $P(H) = 0.541$ (95% HDI [0.28, 0.82]), the highest score of any site at any non-Future epoch, driven by the highest VIMS organic signature of any polar lake combined with confirmed surface liquid and active methane cycling. The smaller polar lacus — Freeman, Oib, Koitere — accumulate more organic material per pixel than the mare shorelines, as resolved by the VIMS spectral data. The three large maria’s primary shorelines (Kraken S, Ligeia E, Punga) rank outside the top 10, replaced by smaller lacus (Bolsena, #9; Cardiel, #10) that have featured in the top 10 at every epoch, driven by concentrated VIMS organic signatures. A 50% reduction in the liquid-hydrocarbon weight ($w_1 : 0.25 \rightarrow 0.13$) reduces Freeman to $P(H) \approx 0.52$, still the highest-scoring site — confirming that the result is driven by the coincidence of multiple features, not by the specific weight of f_1 (Section 4.5). Ontario Lacus (#4, 72°S) scores below the north polar sites primarily because the methane cycle is suppressed at the south pole during the Cassini epoch’s northern summer ($f_4\approx0.45$ vs 0.62 at Freeman) and the lake is actively shrinking. The geographic spread of the top-10 around the pole — spanning 12°W to 327°W longitude (nine sites at 70–83°N, plus Ontario at 72°S), nearly the full polar circle — confirms that the result is robust to specific shoreline geometry. The *Huygens* landing site scores $P(H) = 0.260$ (below the present-epoch median), consistent with its equatorial plains setting lacking both surface liquid and high geodiversity. Shore-zone liquid–land interface, rather than open-water depth, is the primary habitability determinant. These shoreline environments are also the most plausible setting for non-aqueous abiogenesis (Section 1.3): liquid methane solvent, organic sediment from the adjacent terrain, acetylene energy, and the physical gradients created by wave action and seasonal lake-level cycling are all simultaneously present at the shore zone.

The score spread (0.148) is the widest of any epoch except the *Future*, reflecting the strong discriminating power of the VIMS-resolved organic abundance feature combined with the liquid-hydrocarbon signal.

A caveat on the site rankings is warranted. With 95% HDI widths of ~ 0.5 at every pixel (e.g. Freeman: [0.28, 0.82]; Cardiel: [0.09, 0.69]), the credible intervals of rank-1 and rank-10 sites overlap extensively. Individual site rankings within the top 10 are therefore not statistically distinguishable at 95% confidence. The robust finding is the *class* result: polar lake shores as a group score systematically higher than equatorial terrain (group means ~ 0.44 vs ~ 0.25 , a ratio of ~ 1.8), and this separation is stable across all 27 sensitivity combinations. Claims about the specific rank ordering within the top 10 should be understood as maximum-likelihood point estimates rather than statistically significant distinctions.

A hypothetical present-day organism at Freeman Lacus would consume atmospheric acetylene and hydrogen, exhale methane, and wrap itself in an azotosome membrane —

a cell wall made of nitrogen rather than phospholipid (Stevenson et al., 2015). It would experience perpetual darkness for half the Saturnian year, survive methane rainstorms, and regard -179°C as a balmy afternoon. In terrestrial terms, it would be the ultimate cold-adapted methanogen: an organism for which our most extreme Antarctic endolith (Friedmann, 1982) would be, by comparison, basking on a tropical beach.

Rank	Location	Lon ($^{\circ}\text{W}$)	Lat ($^{\circ}\text{N}$)	$P(H)$	Dominant features
1	$\triangle$ Freeman Lacus	210.0	+83.0	0.541	f_2 peak (VIMS), $f_1=0.80$, $f_4=0.62$
2	$\triangle$ Koitere Lacus	154.0	+73.0	0.509	f_2 high (VIMS), $f_1=0.72$, $f_4=0.58$
3	$\triangle$ Oib Lacus	220.0	+82.0	0.509	f_2 high (VIMS), $f_1=0.78$, $f_4=0.60$
4	$\triangle$ Ontario Lacus	179.0	-72.0	0.421	$f_1=0.85$, $f_4=0.45$; south polar
5	$\triangle$ Hammar Lacus	93.5	+70.7	0.409	$f_1=0.82$, $f_4=0.62$
6	$\triangle$ Mackay Lacus	262.8	+77.5	0.408	$f_1=0.80$, $f_4=0.60$
7	$\triangle$ Neagh Lacus	327.0	+72.0	0.404	$f_1=0.78$, $f_4=0.60$
8	$\triangle$ Paxsi Lacus	243.0	+72.0	0.403	$f_1=0.80$, $f_4=0.60$
9	$\triangle$ Bolsena Lacus	12.3	+75.4	0.394	$f_1=0.78$, $f_4=0.60$
10	$\triangle$ Cardiel Lacus	128.0	+78.0	0.393	$f_1=0.78$, $f_4=0.60$
	★ Huygens Site	192.3	-10.6	0.260	Equatorial plains; no liquid
	★ Selk Crater	199.0	+7.0	0.243	$f_7=0.659$; $f_1\approx 0$, $f_4\approx 0$

Table 3.3: Top-10 habitable locations at the *Present* epoch (Cassini era). $\triangle$ = lake/sea shore; ★ = lander site (below line). $P(H)$ from Eq. 2.11.




Figure 3.5: Present-epoch posterior means with 95% HDI for twelve representative sites. Dark blue = lake/sea shore; dark amber = land. Dashed line = prior mean $\mu_0 = 0.331$. HDI widths are similar (~ 0.4 – 0.6), reflecting weak concentration $\kappa = 5$.

3.4.3 Feature Importances

Mean-replacement feature importances from a calibrated Gaussian Naive Bayes classifier trained on posterior-derived labels reveal which features provide the strongest *spatial discrimination* between high- and low-scoring pixels — a distinct quantity from the assigned astrobiological weights w_i . Organic abundance (f_2 , importance 0.399) and methane cycle (f_4 , importance 0.370) dominate the classifier’s discrimination, while liquid hydrocarbon (f_1) drops to near-zero importance despite carrying the highest weight ($w_1 = 0.25$). This apparent discrepancy has a clear physical explanation: f_1 is effectively binary (1.0 at confirmed Birch lake pixels, ≤ 0.05 elsewhere), so it sharply separates $\sim 2\%$ of the globe from the rest but provides no discrimination among the remaining 98%. By contrast, f_2 varies continuously across the globe ($SD \approx 0.27$ from the VIMS spectral mosaic), giving the classifier a rich gradient on which to separate habitable from non-habitable sites. The weights w_i encode which conditions are astrobiologically necessary; the importances reveal which features carry the most spatial information — the two need not agree, and in this framework they do not (Table 3.4).

Feature	w_i	Importance	Ratio
Organic abundance	0.20	0.399	1.99
Methane cycle	0.15	0.370	2.46
Chemical energy	0.20	0.232	1.16
Liquid hydrocarbon	0.25	0.000	0.00
Surface–atmosphere	0.08	0.000	0.00
Topographic complexity	0.06	0.000	0.00
Geomorphological diversity	0.04	0.000	0.00
Subsurface ocean proximity	0.02	0.001	0.05

Table 3.4: Present-epoch mean-replacement feature importances versus prior weights. Under the VIMS-blended organic scoring, f_2 and f_4 provide the strongest geographic discrimination.

3.4.4 Selk Crater

Selk crater (199°W, 7°N; ≈ 80 km; [Hedgepeth et al. 2020](#)) sits at the geological boundary between the Shangri-La dune field (west) and the Adiri bright plains (east). Its SAR-bright annulus is interpreted as the crystallised front of an ancient water–ammonia impact melt ([Wood et al., 2010](#)). VIMS resolves spectral mixing between the tholin signature of Shangri-La and the water-ice signature of excavated ejecta.

The most diagnostically important feature value is $f_7 = 0.659$ ($7.2\times$ the global median of 0.092): the highest geodiversity of any equatorial location, encoding four terrain classes (dune, plains, labyrinth, crater floor) within 45 km. Most other features are near or below their global medians: $f_1 = 0.010$ (no liquid), $f_4 = 0.026$ (equatorial, $23\times$ below median), $f_2 = 0.168$ (crater-floor terrain, below median). The subsurface ocean proximity ($f_8 = 0.030$) is at the global floor (0.030): no SAR-bright annulus is detected at this location.

The aggregate present-epoch posterior is $P(H) = 0.243$ (95% HDI: [0.04, 0.49]) — close to the prior mean, reflecting a site highly distinctive in geomorphological diversity but unremarkable in methane cycle and surface–atmosphere features. Note that applying the feature medians in Table 3.5 directly to Eq. 2.11 gives a slightly lower value ($P(H) \approx 0.215$); the stated 0.243 is the mean posterior over the 5×5 pixel neighbourhood centred on Selk, which samples higher-organic pixels adjacent to the crater annulus. The past-epoch posterior ($P(H) = 0.258$) reflects the declining impact-melt contribution. This terrain-boundary diversity is the Cassini proxy for the organic–water-ice interface targeted by *Dragonfly* (Appendix I): the planned traverse from Belet dunes to the crater floor follows an increasing geodiversity gradient from $f_7 \approx 0.09$ (homogeneous dunes) to $f_7 = 0.659$ (the organic–water-ice contact zone).

Feature	Selk	Global median	Interpretation
f_1 (liquid HC)	0.010	0.050	No present surface liquid; well below global median
f_2 (organic abund.)	0.168	0.510	Below median; crater-floor terrain class
f_3 (chemical energy)	0.232	0.415	Below median; reduced SAR backscatter
f_4 (methane cycle)	0.026	0.600	23 $\times$ below median; equatorial, no lakes
f_5 (surface atm.)	0.010	0.021	Below median
f_6 (topo. complexity)	0.055	0.118	Below median; smooth crater floor
f_7 (geomorph. div.)	0.659	0.092	7.2$\times$ global median ; four terrain classes
f_8 (subsurface ocean)	0.030	0.030	At global floor; no SAR-bright annulus detected

Table 3.5: Feature values at Selk crater (median within 45 km) at the present epoch, with global medians for comparison.




Figure 3.6: Feature profiles and posterior habitability for ten key sites. *Left:* heatmap of normalised feature values. *Right:* posterior mean $P(H | \mathbf{f})$; dashed line = prior mean. Selk’s exceptional geodiversity ($f_7=0.659$) is clearly visible despite its below-average aggregate score.

3.5 Near Future Epoch (+0.25 Gya): The Stable Plateau

The *Near Future* epoch (+0.25 Gya) confirms that Titan’s lake–organic habitability system is robust to small luminosity perturbations of main-sequence solar evolution. Solar luminosity is elevated above present by $\sim 1.4\text{--}6\%$ over the modelled plateau window (Bahcall et al., 2001), corresponding to a surface temperature rise of $\sim 1\text{--}5\text{ K}$ — entirely absorbed by the model and producing no spatially structured change in any feature. The polar seas persist, the methane cycle continues, and the organic inventory grows at the same rate. The spatial correlation between the *Near Future* and *Present* posteriors exceeds $r = 0.999$; the *Near Future* and *Present* maps are functionally equivalent habitability states. All rankings are identical (Table 3.6).

This stability will persist until approximately +4.0 Gya, when increasing solar warming begins to evaporate the polar seas, initiating the transition towards the solvent-free minimum

at $\sim +5.0$ Gya visible in the temporal trend (Figure 3.1).

During this time, any organism that had adapted to the previous Present epoch would barely notice the transition: the ~ 2 K warming amounts to upgrading from a parka to a slightly thinner parka. The methane rain continues, the acetylene supply is uninterrupted, and natural selection has 2.5×10^8 yr of evolutionary stability in which to optimise — a luxury that Earth’s biosphere, punctuated by mass extinctions every $\sim 10^8$ yr, has never enjoyed.




Figure 3.7: *Near Future* epoch (+1.2 Gya) posterior mean habitability. Indistinguishable from Figure 3.4.

Rank	Location	Lon (°W)	Lat (°N)	$P(H)$	Notes
1	△ Freeman Lacus	210.0	+83.0	0.540	Unchanged
2	△ Koitere Lacus	154.0	+73.0	0.509	Unchanged
3	△ Oib Lacus	220.0	+82.0	0.508	Unchanged
4	△ Ontario Lacus	179.0	−72.0	0.420	Unchanged
5	△ Hammar Lacus	93.5	+70.7	0.408	Unchanged
6	△ Mackay Lacus	262.8	+77.5	0.407	Unchanged
7	△ Neagh Lacus	327.0	+72.0	0.403	Unchanged
8	△ Paxsi Lacus	243.0	+72.0	0.402	Unchanged
9	△ Bolsena Lacus	12.3	+75.4	0.393	Unchanged
10	△ Cardiel Lacus	128.0	+78.0	0.392	Unchanged
	★ Huygens Site	192.3	−10.6	0.259	Unchanged
	★ Selk Crater	199.0	+7.0	0.242	Unchanged

Table 3.6: *Near Future* (+0.25–+1.2 Gya) *top*–10 : *effectively identical to Present across the entire stable plateau*. ★ = *lander site* (below line).

3.6 *Future* Epoch (+5.9 Gya): The Red-Giant Ocean

The *Future* epoch (+5.9 Gya) models the most dramatic transformation of Titan’s habitability landscape. As the Sun evolves off the main sequence, its luminosity increases by orders of magnitude, raising Titan’s surface temperature well above the water–ammonia eutectic (176 K; [Grasset and Pargamin 2005](#)). At +5.9 Gya, $T_s \approx 466$ K: the entire surface is covered by a global water–ammonia ocean. The polar hydrocarbon seas, which dominated the *Present* and *Near Future* rankings, evaporated by $\sim +4.5$ Gya. In their place, ~ 10 Gyr of accumulated tholins — previously inert at 94 K — dissolve into warm water–ammonia, creating what [Lorenz et al. \(1997\)](#) described as a planetary-scale prebiotic ocean. The model applies $f_1 = f_8 = 1.0$ everywhere, with organic accumulation at $2.5\times$ present. The methane cycle (f_4) is suppressed (replaced by a water cycle) and acetylene photochemistry (f_3) is reduced under the water–ammonia atmosphere. Between $\sim +4.0$ and +5.1 Gya, Titan is transiently solvent-free — a ~ 100 Myr astrobiological desert when neither hydrocarbon nor water-based habitability is available.

The *Future* map (Figure 3.8) shows elevated posteriors across the entire surface — the highest global mean of any epoch. The sharp polar/equatorial contrast of the *Present* epoch is reduced but not eliminated: former polar lake basins, whose VIMS-derived organic abundance is amplified $2.5\times$ by the temporal organic accumulation factor, score substantially higher than inundated equatorial terrain. Comparing with the temporal trend (Figure 3.1), this epoch represents the recovery from the solvent-free minimum at $\sim +5.0$ Gya; the north polar and equatorial curves approach one another but do not equalise, because the organic accumulation advantage of the polar basins persists through the ocean-covered epoch. The ranking is dominated by former polar lake sites (Table 3.7). Mackay Lacus ($P(H)=0.763$) leads at this epoch, overtaking Freeman (#2, $P(H)=0.723$), because the $2.5\times$ organic amplification at the Future epoch raises Mackay’s VIMS signal above Freeman’s — a reversal of the Present ranking driven by Mackay’s slightly higher pre-amplification organic score at the specific pixel sampled. Both scores exceed the theoretical bound of $P_{\max} = 0.696$ (which

assumes all features ≤ 1.0) because the temporal scaling pushes the effective organic feature above unity. Former polar lake basins occupy all ten top-10 positions, with scores ranging from 0.641 to 0.763 — far above the prior-epoch maximum of 0.541 (Freeman, Present). Cryovolcanic sites (Hotei, Sotra) and inundated equatorial dune fields (Belet, Shangri-La) rank below the top 10.

The score spread (0.122: from 0.641 to 0.763) is driven entirely by the VIMS-derived organic abundance variation amplified by the $2.5\times$ temporal factor — the remaining features (f_5 , f_6) contribute negligible differentiation with $f_1=1.0$ everywhere. The red-giant ocean window (+5.1 to +6.0 Gya, ~ 0.9 Gyr) is the longest continuous period of global liquid coverage in Titan’s modelled history, combining maximal organic inventory with warm water–ammonia solvent — an extraordinary future habitability scenario unmatched by any other epoch or body in the post-main-sequence solar system. Of all epochs modelled, this represents the strongest candidate for aqueous abiogenesis: warm liquid water, the full accumulated organic inventory of ~ 10 Gyr of photochemistry, and a timescale ($\sim 10^8$ yr) comparable to the interval between the formation of Earth’s oceans and the earliest evidence of terrestrial life.

An organism that had spent billions of years perfecting cryogenic methane biochemistry would face the most dramatic environmental challenge in Titan’s history — analogous to, but far more extreme than, the terrestrial Great Oxidation Event. If the transition through the solvent-free desert at $\sim +5.0$ Gya does not sterilise the surface entirely, the survivors would find themselves in warm water with an effectively unlimited organic buffet. With ~ 0.9 Gyr of sustained molecular evolution in conditions comparable to those that produced the first eukaryotes on Earth (Knoll, 2014), the Future ocean is — at least in principle — the epoch where Titan stops being a world of patient, diffusing chemotrophs and starts being a world where biology could get ambitious.




Figure 3.8: Future epoch (+5.9Gya; red-giant ocean) posterior mean habitability. f_1 and f_8 overridden to 1.0 globally. Spatial variation reflects f_2, f_5, f_6 differences. Compare with Figure 3.4.

Rank	Location	Lon (°W)	Lat (°N)	$P(H)$	Dominant features
1	△ Mackay (ocean floor)	262.8	+77.5	0.763	$f_1=1.0, f_2 \text{ peak} \times 2.5$
2	△ Freeman (ocean floor)	210.0	+83.0	0.723	$f_1=1.0, f_2 \text{ high} \times 2.5$
3	△ Hammar (ocean floor)	93.5	+70.7	0.708	$f_1=1.0, f_2 \text{ high} \times 2.5$
4	△ Uvs (ocean floor)	262.0	+74.5	0.700	$f_1=1.0, f_2 \text{ high} \times 2.5$
5	△ Waikare (ocean floor)	185.0	+75.0	0.682	$f_1=1.0, f_2 \text{ elevated} \times 2.5$
6	△ Ontario (ocean floor)	179.0	-72.0	0.679	$f_1=1.0$; south polar f_2
7	△ Logtak (ocean floor)	197.5	+74.0	0.675	$f_1=1.0, f_2 \text{ high} \times 2.5$
8	△ Paxsi (ocean floor)	243.0	+72.0	0.671	$f_1=1.0$, organics dissolved
9	△ Oib (ocean floor)	220.0	+82.0	0.641	$f_1=1.0$, organics dissolved
10	△ Cardiel (ocean floor)	128.0	+78.0	0.641	$f_1=1.0$, organics dissolved
	★ Selk Crater	199.0	+7.0	0.419	$f_1=1.0$; low f_2 limits rank
	★ Huygens Site	192.3	-10.6	0.409	$f_1=1.0$; submerged

Table 3.7: Top-10 habitable locations at the *Future* epoch (+5.9 Gya). △ = former lake (now ocean floor); ★ = lander site (below line: outside top 10). $P(H)$ from Eq. 2.11.

4 Discussion

4.1 Dragonfly Landing Site Validation

The pipeline independently validates the *Dragonfly* mission’s selection of Selk crater, but through a different route than a global ranking. The maps in Chapter 3 quantify *present-day solvent* habitability; *Dragonfly* targets *past liquid-water* prebiotic chemistry. These are two scientifically distinct frameworks that correctly produce different site rankings: a model that captures present-day solvent habitability *should* rank polar lake shores above a dry equatorial crater (Appendix I). Selk’s present-epoch $P(H) = 0.243$ is therefore not a contradiction of the mission rationale but a consequence of targeting different habitability.

The independent validation comes from the geomorphological diversity feature: Selk’s $f_7 = 0.659$ ($7.2\times$ the global median of 0.092) is derived from Cassini SAR data alone, without reference to mission planning, and without any weight or feature specification adjusted after examining Selk. This terrain-boundary diversity — four terrain classes within 45 km — identifies the organic–water-ice contact zone that is the mission’s primary scientific target.

4.2 Titan in Solar System Context

Present-epoch Titan is distinguished from all other solar system candidates by combining a confirmed cycling surface solvent with a large organic inventory. Future Titan (+5.9 Gya) uniquely pairs a water-based ocean with ~ 10 Gyr of accumulated organics — the most favourable habitability combination modelled for any body in the post-main-sequence solar system.

For context, applying the framework’s logic to the closest competing candidates yields instructive contrasts. Enceladus possesses subsurface liquid water and hydrothermal activity (Waite et al., 2017), but its surface is essentially solvent-free (no liquid layer, $f_1 \approx 0$) and its organic inventory is limited to what is ejected through the plumes; an analogue framework would place Enceladus substantially below the present-epoch Titan polar lakes in surface habitability, though above Titan in subsurface habitability. Europa’s subsurface ocean is similarly promising (Hendrix et al., 2019) but is inaccessible at the surface. By the metric this thesis quantifies — co-occurrence of surface solvent, organic substrate, chemical energy, and solvent cycling — present-epoch Titan has no peer among currently accessible solar system bodies. Affholder et al. (2021) estimated the probability of methanogenesis on Enceladus at a level consistent with existing data. For Titan’s own subsurface, Affholder et al. (2025) used bioenergetic modelling to find that the total viable biomass in any water-bearing interior layer would amount to only a few grams to a few kilograms of carbon — fewer than one microbial cell per kilogram of water throughout the vast ocean that was previously assumed (Neish et al., 2024b). The present framework’s $P(H) = 0.541$ for Freeman Lacus does not directly translate to a probability of life, but it does identify

Titan’s polar lake shores as providing the most co-located habitability preconditions of any accessible solar system surface — a distinct and complementary result from the subsurface biomass estimates.

4.3 Implications for Abiogenesis

The habitability maps produced in this thesis identify where the preconditions for life are co-located, but the temporal reconstruction also reveals which epochs and sites are most relevant to the *origin* of life. Two epochs stand out.

The *Past* epoch (−3.5 Gya) offers the most direct route to aqueous abiogenesis through impact-melt ponds at crater sites (Section 3.2). Liquid water, accumulated tholin substrate, atmospheric HCN, and ammonia are simultaneously present for 10^3 – 10^4 yr per impact — long enough for Strecker amino acid synthesis (Neish et al., 2010) but possibly too short for the transition from prebiotic chemistry to self-replicating systems. The model ranks these sites 11th and below in solvent habitability, precisely because the habitability proxy captures present-day solvent conditions rather than transient aqueous chemistry — an important caveat for interpreting the maps.

The *Future* epoch (+5.9 Gya) represents the strongest abiogenesis scenario in the entire reconstruction. The global water–ammonia ocean dissolves ~ 10 Gyr of accumulated tholins, providing sustained contact between warm water solvent, complex organic substrate, and solar energy over ~ 0.9 Gyr — conditions broadly analogous to those invoked for early Earth but with a far richer initial organic inventory. Unlike the Archean Earth, where CO_2 and CH_4 greenhouse gases sustained liquid water under a faint Sun (Catling and Zahnle, 2020), the Future Titan ocean is maintained by direct stellar luminosity increase — a fundamentally different atmospheric context that may favour different prebiotic pathways. Unlike the transient impact-melt ponds, the timescale is comfortably long for molecular evolution: the ~ 0.9 Gyr ocean window is comparable to the interval between the formation of Earth’s oceans and the earliest evidence of terrestrial life at ~ 3.5 Ga (Brasier et al., 2015). Former polar lake-basin sites lead the ranking at this epoch, with Mackay Lacus ($P(H)=0.763$) overtaking Freeman ($P(H)=0.723$) under the $2.5\times$ temporal organic amplification, producing the highest scores of any site at any epoch.

The present-day polar lake shores, despite scoring highest for solvent habitability, present the most speculative abiogenesis scenario: non-aqueous abiogenesis in liquid methane at 94 K, requiring compartmentalisation via azotosome membranes (Stevenson et al., 2015) and energy from acetylenotrophy (McKay and Smith, 2005). No terrestrial analogue exists for the origin of life in a non-aqueous medium, and the low temperature severely constrains reaction kinetics. It is important to note that this framework quantifies habitability conditions — the environmental prerequisites for life — rather than biosignatures in the sense of Schwieterman and Leung (2024), which would require in-situ atmospheric or surface measurements from *Dragonfly* or a comparable mission. The habitability maps therefore highlight a fundamental tension: the sites with the highest present-day habitability scores are the least well understood in terms of abiogenesis potential, while the sites with the strongest abiogenesis credentials (impact craters, the future ocean) score modestly or are

temporally inaccessible.

4.4 Limitations of the Framework

The framework relies on several simplifying assumptions, documented in full in Appendix E. The most consequential are summarised here.

Conditional feature independence. The Beta-conjugate model treats the eight features as contributing independently. In reality, f_1 (liquid hydrocarbon), f_4 (methane cycle), and f_5 (surface–atmosphere interaction) are physically coupled at lake margins: all three are largely driven by whether surface liquid is present at a pixel. Collectively they carry weight $w_1 + w_4 + w_5 = 0.48$ — nearly half the total — so lake-shore sites receive substantially correlated evidence for the same underlying condition. If this correlated block were collapsed to a single effective feature with weight ~ 0.25 , Freeman’s score would reduce by approximately 0.05–0.08 (from 0.541 to ~ 0.46 –0.49), and the top-10 range would compress by a similar amount. The qualitative result is unchanged: polar lake shores remain the highest-scoring class and the *Dragonfly* validation through f_7 is entirely unaffected. Future work should replace the independence assumption with a full covariance structure estimated from high-resolution SAR patches (Section 4.6). The sensitivity analysis (Section 4.5) additionally bounds variation due to prior specification: $|\Delta P(H)| \leq 0.06$ for sites near the prior mean, comparable to the prior-specification uncertainty.

Static geomorphological template. Feature maps are derived from 2004–2017 Cassini observations and applied uniformly to all non-present epochs via scalar scale functions. Real geological changes over billions of years — dune migration, basin infilling, tectonic modification — are unrepresented. This approximation is justified for the broad-scale habitability classification pursued here, but limits the resolution of transitional epochs.

Past-epoch Cassini contamination. The geomorphological map reflects the present surface; at the LHB, terrain distributions were different. Scale functions partially correct for this by reducing organic abundance, but the spatial pattern of f_7 is held fixed.

4.5 Sensitivity of Results to Prior Specification

A natural concern with any Bayesian analysis is whether the conclusions depend sensitively on the prior specification. Three parameters govern the prior: κ (concentration), λ (likelihood sharpness), and the feature weights w_i . Figure 4.1 shows the effect of varying each at three representative sites: Freeman (rank 1, high organic + liquid), Belet (equatorial organic site, outside top 10), and Selk (below prior mean, high geodiversity).

The key finding is robustness. Freeman retains rank 1 in all $3 \times 3 \times 3 = 27$ parameter combinations. Halving κ from 5 to 2.5 increases Freeman’s $P(H)$ by +0.065 and reduces Selk’s by -0.020 ; doubling to 10 reverses these changes (-0.068 and $+0.022$). For sites near the prior mean (Belet, Selk), the maximum shift is $|\Delta P(H)| \leq 0.06$ — well below the 95% HDI width of ≈ 0.5 at any single site. The qualitative ordering (organic-rich polar lake shores $>$ other lacus $>$ equatorial sites) is preserved in all 27 combinations. Crucially, reducing w_1 from 0.25 to 0.13 (halving the weight of liquid hydrocarbon) only reduces Freeman’s $P(H)$

from 0.541 to ~ 0.52 — Freeman remains the highest-scoring site because its VIMS-derived organic abundance, not the liquid-hydrocarbon feature alone, drives the ranking.



Figure 4.1: Sensitivity of $P(H | \mathbf{f})$ to prior specification. Each row varies one parameter; each column shows one site. Deviations ≤ 0.06 for near-prior-mean sites; qualitative rankings preserved across all 27 parameter combinations.

4.6 Future Work

The three most impactful extensions are: (1) post-*Dragonfly* recalibration of f_2 and f_8 weights using DraMS and DraGNS measurements at Selk; (2) replacing the independence assumption with a full covariance structure estimated from high-resolution SAR patches; and (3) additional temporal anchors at -2.0 to -0.5 Gya where the liquid-hydrocarbon ramp most sensitively determines the magnitude and character of the polar lead during the transition from organic-dominated to lake-dominated habitability.

4.7 Additional Datasets

Two recently released Cassini-derived datasets from the Cornell Cassini RADAR team archive (<https://data.astro.cornell.edu>) could improve the spatial coverage of features f_3 and f_5 in future work:

Incidence-angle-corrected SAR mosaics. Birch et al. produced four equatorial SAR mosaic tiles (60°S–60°N, 256 ppd) with empirical σ_0 incidence-angle correction, used as basemaps in several geomorphological mapping campaigns (Birch et al., 2025; Malaska et al., 2025). The standard USGS global SAR mosaic (c1on180 product) does not apply uniform incidence-angle normalisation across flyby swaths, producing the intensity discontinuity at $\sim 180^\circ\text{W}$ visible in the f_3 acetylene energy feature (Figure G.1; equatorial step $|\Delta f_3| \approx 0.12\text{--}0.15$). If the Birch σ_0 -corrected tiles reduce this step at the 180°W boundary, replacing the current SAR backscatter layer in f_3 with the corrected mosaic would eliminate the largest single spatial artefact in the feature set.

Extended fluvial channel mapping. The f_5 surface–atmosphere interaction feature includes a 30% channel-density component derived from the Miller et al. (2021) global channel map ($\sim 4\,000$ channel segments, 1.2 MB shapefile). An extended channel mapping dataset by Verdi (2025, unpublished; Cornell data archive, 8.8 GB) may provide substantially greater spatial coverage of fluvial features mapped from SAR imagery, potentially reducing the sparseness of f_5 at mid-latitudes where Miller’s mapping was limited by SAR swath availability. Integration would require validation of the mapping methodology and comparison of spatial coverage with the existing Miller et al. product.

4.8 Conclusions

This thesis presents the first quantitative, spatially-resolved, multi-epoch Bayesian habitability map of Titan’s surface, derived from eight Cassini observational proxy features. Conclusions are organised by research question.

RQ1: Present-epoch spatial distribution. North polar lake margins are the most habitable present-day environments. Freeman Lacus achieves $P(H) = 0.541$ (rank 1), driven by the highest VIMS-derived organic abundance of any polar lake combined with confirmed surface liquid, active methane cycling, and elevated geodiversity. The three large maria’s primary shorelines (Kraken, Ligeia, Punga) rank outside the top 10 at all epochs, replaced by smaller organic-rich lacus (Bolsena, Cardiel) at positions 9–10. All ten top-10 sites are lake or lacus margins at every epoch.

RQ2: LHB habitability. Polar proto-basin sites lead the *Past* ranking (Freeman $P(H)=0.464$) because the VIMS-derived organic abundance at these locations dominates the scoring even before lakes have formed. Cryovolcanic sites and craters rank outside the top 10 in the 8-feature formula; their astrobiological significance lies in providing transient liquid water for Strecker synthesis, a distinct habitability metric not captured by the solvent-habitability model.

RQ3: Temporal trajectory. The trajectory is non-monotonic across three regimes. (i) LHB (–3.8 to –1.5 Gya): organic-accumulation dominance, where polar proto-basins

lead on VIMS-derived f_2 alone; scores are compressed because the liquid-hydrocarbon signal is absent. (ii) Polar-lake dominance (-1.0 to $+4.0$ Gya): the liquid-hydrocarbon feature amplifies the existing polar lead, producing the maximum spatial contrast; lake or lacus margins occupy all top-10 positions across three epochs with a stable composition (Freeman #1 throughout); the present-day habitability structure is a plateau extending at least 250 Myr into the future. (iii) Global ocean ($+5.1$ to $+6.0$ Gya): former polar lake basins lead (Mackay $P(H)=0.763$, Freeman $P(H)=0.723$) as $2.5\times$ temporal organic scaling amplifies the VIMS organic signal; the theoretical $P_{\max}$ of 0.696 is exceeded. A transient solvent-free minimum at $\sim+5.0$ Gya separates regimes (ii) and (iii). The interval between $+4.0$ and $+5.1$ Gya — roughly 1.1 Gyr — represents an astrobiological discontinuity of considerable interest: the polar seas have evaporated but the water–ammonia eutectic has not yet been crossed, leaving Titan’s surface entirely without liquid solvent. This raises a question the present framework cannot answer: whether life adapted to the methane-lake regime of regime (ii) could survive or leave biosignatures through this desert epoch and seed the red-giant ocean of regime (iii). Dormancy strategies analogous to terrestrial endospore formation, or subsurface refuge in deep ice or the subsurface ocean, represent plausible but unmodelled pathways. The 1.1 Gyr gap is arguably the most important unresolved question raised by this temporal reconstruction.

RQ4: Dragonfly validation. Selk crater has the highest geodiversity of any equatorial site ($f_7=0.659$, $7.2\times$ the global median), derived independently from SAR data. Its lower present-epoch score ($P(H)=0.243$) is consistent: *Dragonfly* targets past water–organic chemistry, not present-day solvent habitability (Appendix I).

Together, these results demonstrate that Titan possesses a multilayered astrobiological significance from its geological past through the present to the far future, quantifiable from the Cassini dataset with sufficient precision to guide mission planning and interpret in-situ measurements. The temporal reconstruction reveals a fundamental tension between habitability and abiogenesis (Section 4.3): the sites with the highest present-day solvent habitability scores (polar lake shores) are the least understood in terms of origin-of-life potential, while the environments with the strongest abiogenesis credentials — LHB impact-melt ponds and the red-giant ocean — score modestly or lie in Titan’s temporal past and future respectively.

Bibliography

- Affholder, A., Guyot, F., Sauterey, B., Ferrière, R., and Mazevet, S. (2021). Bayesian analysis of Enceladus’s plume data to assess methanogenesis. *Nature Astronomy*, 5:805–814. doi:[10.1038/s41550-021-01372-6](https://doi.org/10.1038/s41550-021-01372-6).
- Affholder, A., Higgins, P. M., Nixon, C. A., Melwani Daswani, M., Miller, K. E., and Sotin, C. (2025). The viability of glycine fermentation in Titan’s subsurface ocean. *The Planetary Science Journal*, 6(4):86. doi:[10.3847/PSJ/adbc66](https://doi.org/10.3847/PSJ/adbc66).
- Aharonson, O., Hayes, A. G., Lunine, J. I., Lorenz, R. D., Allison, M. D., and Elachi, C. (2009). An asymmetric distribution of lakes on Titan as a possible consequence of orbital forcing. *Nature Geoscience*, 2:851–854. doi:[10.1038/ngeo698](https://doi.org/10.1038/ngeo698).
- Artemieva, N. and Lunine, J. I. (2003). Cratering on Titan: impact melt, ejecta, and the fate of surface organics. *Icarus*, 164(2):471–480. doi:[10.1016/S0019-1035\(03\)00148-9](https://doi.org/10.1016/S0019-1035(03)00148-9).
- Bahcall, J. N., Pinsonneault, M. H., and Basu, S. (2001). Solar models: Current epoch and time dependences, neutrinos, and helioseismological properties. *The Astrophysical Journal*, 555:990–1012. doi:[10.1086/321493](https://doi.org/10.1086/321493).
- Barnes, J. W. et al. (2007). Global-scale surface spectral variations on Titan seen from Cassini/VIMS. *Icarus*, 186(1):242–258. doi:[10.1016/j.icarus.2006.08.021](https://doi.org/10.1016/j.icarus.2006.08.021).
- Benner, S. A., Ricardo, A., and Carrigan, M. A. (2004). Is there a common chemical model for life in the universe? *Current Opinion in Chemical Biology*, 8:672–689. doi:[10.1016/j.cbpa.2004.10.003](https://doi.org/10.1016/j.cbpa.2004.10.003).
- Birch, S. P. D. et al. (2025). Titan’s volatile cycle. In *Titan After Cassini-Huygens*, chapter 10. Cambridge University Press.
- Brasier, M. D., Antcliffe, J., Saunders, M., and Wacey, D. (2015). Changing the picture of Earth’s earliest fossils (3.5–1.9 Ga) with new approaches and new discoveries. *Proceedings of the National Academy of Sciences*, 112(16):4859–4864. doi:[10.1073/pnas.1405338111](https://doi.org/10.1073/pnas.1405338111).
- Brown, R. H. et al. (2008). The identification of liquid ethane in Titan’s Ontario Lacus. *Nature*, 454:607–610. doi:[10.1038/nature07100](https://doi.org/10.1038/nature07100).
- Cable, M. L., Hörst, S. M., Hodyss, R., Beauchamp, P. M., Smith, M. A., and Willis, P. A. (2012). Titan tholins: simulating Titan organic chemistry in the Cassini-Huygens era. *Chemical Reviews*, 112:1882–1909. doi:[10.1021/cr200221x](https://doi.org/10.1021/cr200221x).
- Castlevetro, L. (2026). Title of the thesis. Master’s thesis, University of Cambridge.
- Catling, D. C., Krissansen-Totton, J., Kiang, N. Y., et al. (2018). Exoplanet biosignatures: A framework for their assessment. *Astrobiology*, 18(6):709–738. doi:[10.1089/ast.2017.1737](https://doi.org/10.1089/ast.2017.1737).

- Catling, D. C. and Zahnle, K. J. (2020). The Archean atmosphere. *Science Advances*, 6(9):eaax1420. doi:[10.1126/sciadv.eaax1420](https://doi.org/10.1126/sciadv.eaax1420).
- Cockell, C. S., Bush, T., Bryce, C., Direito, S., Fox-Powell, M., Harrison, J. P., Lammer, H., Landenmark, H., Martin-Torres, J., Nicholson, N., Noack, L., O'Malley-James, J., Payler, S. J., Rushby, A., Samuels, T., Schwendner, P., Wadsworth, J., and Zorzano, M. P. (2016). Habitability: A review. *Astrobiology*, 16(1):89–117. doi:[10.1089/ast.2015.1295](https://doi.org/10.1089/ast.2015.1295).
- Corlies, P., Hayes, A. G., Birch, S. P. D., Lorenz, R., Stiles, B. W., Kirk, R., and Poggiali, V. (2017). Titan’s topography and shape at the end of the cassini mission. *Geophysical Research Letters*, 44(23):11754–11761. doi:[10.1002/2017GL075518](https://doi.org/10.1002/2017GL075518).
- Crick, C. (2026). Title of the thesis. Master’s thesis, University of Cambridge.
- DeGroot, M. H. and Schervish, M. J. (2012). *Probability and Statistics*. Pearson, Boston, 4th edition.
- Flasar, F. M. et al. (2005). Titan’s atmospheric temperatures, winds, and composition. *Science*, 308:975–978. doi:[10.1126/science.1111150](https://doi.org/10.1126/science.1111150).
- Friedmann, E. I. (1982). Endolithic microorganisms in the Antarctic cold desert. *Science*, 215(4536):1045–1053. doi:[10.1126/science.215.4536.1045](https://doi.org/10.1126/science.215.4536.1045).
- Fulchignoni, M. et al. (2005). In situ measurements of the physical characteristics of Titan’s environment. *Nature*, 438:785–791. doi:[10.1038/nature04314](https://doi.org/10.1038/nature04314).
- Gelman, A., Carlin, J. B., Stern, H. S., Dunson, D. B., Vehtari, A., and Rubin, D. B. (2013). *Bayesian Data Analysis*. Chapman and Hall/CRC, 3rd edition.
- Grasset, O. and Pargamin, J. (2005). The ammonia–water system at high pressures: Implications for the methane of Titan. *Planetary and Space Science*, 53(4):371–384. doi:[10.1016/j.pss.2004.09.042](https://doi.org/10.1016/j.pss.2004.09.042).
- Hayes, A., Aharonson, O., Callahan, P., et al. (2008). Hydrocarbon lakes on Titan: Distribution and interaction with a porous regolith. *Geophysical Research Letters*, 35(9):L09204. doi:[10.1029/2007GL032324](https://doi.org/10.1029/2007GL032324).
- Hayes, A. G. (2016). The lakes and seas of Titan. *Annual Review of Earth and Planetary Sciences*, 44:57–83. doi:[10.1146/annurev-earth-060115-012247](https://doi.org/10.1146/annurev-earth-060115-012247).
- Hedgepeth, J. E., Neish, C. D., Turtle, E. P., and Stiles, B. W. (2020). Titan’s impact crater population after Cassini. *Icarus*, 344:113664. doi:[10.1016/j.icarus.2020.113664](https://doi.org/10.1016/j.icarus.2020.113664).
- Hemingway, D., Nimmo, F., Zebker, H., and Iess, L. (2013). A rigid and weathered ice shell on Titan. *Nature*, 500:550–552. doi:[10.1038/nature12400](https://doi.org/10.1038/nature12400).
- Hendrix, A. R., Hurford, T. A., Barge, L. M., Bland, M. T., Bowman, J. S., Brinckerhoff, W., Buratti, B. J., Cable, M. L., Castillo-Rogez, J., et al. (2019). The NASA roadmap to ocean worlds. *Astrobiology*, 19(1):1–27. doi:[10.1089/ast.2018.1955](https://doi.org/10.1089/ast.2018.1955).

- Iess, L., Jacobson, R. A., Ducci, M., et al. (2012). The tides of Titan. *Science*, 337:457–459. doi:[10.1126/science.1219631](https://doi.org/10.1126/science.1219631).
- Janssen, M. A. et al. (2016). Titan’s surface at 2.18-cm wavelength imaged by the Cassini RADAR radiometer. *Icarus*, 270:443–459. doi:[10.1016/j.icarus.2015.09.027](https://doi.org/10.1016/j.icarus.2015.09.027).
- Jennings, D. E., Tokano, T., Cottini, V., Nixon, C. A., Achterberg, R. K., Flasar, F. M., Kunde, V. G., Romani, P. N., Samuelson, R. E., Segura, M. E., Gorius, N. J. P., Guandique, E., Kaelberer, M. S., and Coustenis, A. (2019). Titan surface temperatures during the Cassini mission. *The Astrophysical Journal Letters*, 877(1):L8. doi:[10.3847/2041-8213/ab1f91](https://doi.org/10.3847/2041-8213/ab1f91).
- Knoll, A. H. (2014). Paleobiological perspectives on early eukaryotic evolution. *Cold Spring Harbor Perspectives in Biology*, 6:a016121. doi:[10.1101/cshperspect.a016121](https://doi.org/10.1101/cshperspect.a016121).
- Lavvas, P., Griffith, C. A., and Yelle, R. V. (2011). Condensation in Titan’s atmosphere at the Huygens landing site. *Icarus*, 215:732–750. doi:[10.1016/j.icarus.2011.06.040](https://doi.org/10.1016/j.icarus.2011.06.040).
- Le Mouélic, S. et al. (2019). The Cassini VIMS archive of Titan: from browse products to global infrared color maps. *Icarus*, 319:121–132. doi:[10.1016/j.icarus.2018.09.017](https://doi.org/10.1016/j.icarus.2018.09.017).
- Little, R. J. A. and Rubin, D. B. (2019). *Statistical Analysis with Missing Data*. John Wiley & Sons, 3rd edition. doi:[10.1002/9781119482260](https://doi.org/10.1002/9781119482260).
- Loison, J.-C. et al. (2015). The neutral photochemistry of nitriles, amines and imines in the atmosphere of Titan. *Icarus*, 247:218–247. doi:[10.1016/j.icarus.2014.09.039](https://doi.org/10.1016/j.icarus.2014.09.039).
- Lopes, R. M. C. et al. (2019a). Titan as revealed by the Cassini radar. *Space Science Reviews*, 215:33. doi:[10.1007/s11214-019-0598-6](https://doi.org/10.1007/s11214-019-0598-6).
- Lopes, R. M. C. et al. (2020). A close look at Titan’s geology. *Journal of Geophysical Research: Planets*. doi:[10.1029/2019JE006239](https://doi.org/10.1029/2019JE006239).
- Lopes, R. M. C., Malaska, M. J., Schoenfeld, A. M., et al. (2019b). A global geomorphologic map of Saturn’s moon Titan. *Nature Astronomy*, 4:228–233. doi:[10.1038/s41550-019-0917-6](https://doi.org/10.1038/s41550-019-0917-6).
- Lorenz, R. D. (2014). The flushing of Ligeia: Composition variations across Titan’s seas in a simple diffusion model. *Geophysical Research Letters*, 41(15):5764–5770. doi:[10.1002/2014GL061133](https://doi.org/10.1002/2014GL061133).
- Lorenz, R. D. et al. (2006). The sand seas of Titan: Cassini RADAR observations of longitudinal dunes. *Science*, 312:724–727. doi:[10.1126/science.1123257](https://doi.org/10.1126/science.1123257).
- Lorenz, R. D. et al. (2008). Titan’s inventory of organic surface materials. *Geophysical Research Letters*, 35:L02206. doi:[10.1029/2007GL032118](https://doi.org/10.1029/2007GL032118).
- Lorenz, R. D. et al. (2018). Dragonfly: A rotorcraft lander concept for scientific exploration at Titan. *Johns Hopkins APL Technical Digest*, 34(3):374–387. URL <https://www.jhuapl.edu/techdigest>.

- Lorenz, R. D., Lunine, J. I., and McKay, C. P. (1997). Titan under a red giant sun: A new kind of “habitable” moon. *Geophysical Research Letters*, 24(22):2905–2908. doi:[10.1029/97GL52843](https://doi.org/10.1029/97GL52843).
- Lorenz, R. D., MacKenzie, S. M., Neish, C. D., Le Gall, A., Turtle, E. P., Barnes, J. W., Trainer, M. G., Werynski, A., Hedgepeth, J. E., and Karkoschka, E. (2021). Selection and characteristics of the Dragonfly landing site near Selk crater, Titan. *The Planetary Science Journal*, 2(1):24. doi:[10.3847/PSJ/abd08d](https://doi.org/10.3847/PSJ/abd08d).
- Lorenz, R. D., Stiles, B. W., Aharonson, O., et al. (2013). A global topographic map of Titan. *Icarus*, 225(1):367–377. doi:[10.1016/j.icarus.2013.04.002](https://doi.org/10.1016/j.icarus.2013.04.002).
- Madan, R. et al. (2026). Amino acid synthesis in Titan impact-melt environments. *Icarus*. In press / preprint; DOI to be confirmed at publication.
- Malaska, M. J. et al. (2020). Labyrinth terrain on Titan. *Icarus*, 344:113764. doi:[10.1016/j.icarus.2020.113764](https://doi.org/10.1016/j.icarus.2020.113764).
- Malaska, M. J. et al. (2025). Global geology of Titan. In *Titan After Cassini-Huygens*, chapter 9. Elsevier. doi:[10.1016/B978-0-323-99161-2.00013-9](https://doi.org/10.1016/B978-0-323-99161-2.00013-9).
- Mastrogiuseppe, M., Hayes, A. G., Poggiali, V., et al. (2019). Deep and methane-rich lakes on Titan. *Nature Astronomy*, 3:535–542. doi:[10.1038/s41550-019-0714-2](https://doi.org/10.1038/s41550-019-0714-2).
- McKay, C. P. (2016). Titan as the abode of life. *Life*, 6(1):8. doi:[10.3390/life6010008](https://doi.org/10.3390/life6010008).
- McKay, C. P. and Smith, H. D. (2005). Possibilities for methanogenic life in liquid methane on the surface of Titan. *Icarus*, 178(1):274–276. doi:[10.1016/j.icarus.2005.05.018](https://doi.org/10.1016/j.icarus.2005.05.018).
- Méndez, A., Rivera-Valentín, E. G., Schulze-Makuch, D., Filiberto, J., Ramírez, R. M., Wood, T. E., et al. (2021). Habitability models for astrobiology. *Astrobiology*, 21(8):1017–1027. doi:[10.1089/ast.2020.2342](https://doi.org/10.1089/ast.2020.2342).
- Meyer-Dombard, D. R. et al. (2025). Habitability of Titan’s interior and surface. In *Titan After Cassini-Huygens*, chapter 14. Elsevier. doi:[10.1016/B978-0-323-99161-2.00003-6](https://doi.org/10.1016/B978-0-323-99161-2.00003-6).
- Miller, J., Hayes, A. G., Corlies, P., Birch, S. P. D., Lorenz, R. D., and Turtle, E. P. (2021). Global fluvial channel map of Titan from Cassini SAR. *Journal of Geophysical Research: Planets*, 126(12). doi:[10.1029/2021JE006955](https://doi.org/10.1029/2021JE006955).
- Mitchell, J. L. and Lora, J. M. (2016). The climate of Titan. *Annual Review of Earth and Planetary Sciences*, 44:353–380. doi:[10.1146/annurev-earth-060115-012054](https://doi.org/10.1146/annurev-earth-060115-012054).
- Neish, C. D., Blichert-Toft, J., et al. (2024a). Evidence for ancient subsurface ocean–surface exchange on Titan from crater ejecta composition. *Nature Geoscience*, 17:238–244. doi:[10.1038/s41561-024-01357-4](https://doi.org/10.1038/s41561-024-01357-4).
- Neish, C. D., Lorenz, R. D., Turtle, E. P., Barnes, J. W., Trainer, M. G., Stiles, B., Kirk, R., Hibbitts, C. A., and Malaska, M. J. (2018). Strategies for detecting biological molecules on Titan. *Astrobiology*, 18(5):571–585. doi:[10.1089/ast.2017.1758](https://doi.org/10.1089/ast.2017.1758).

- Neish, C. D., Lorenz, R. D., Turtle, E. P., Barnes, J. W., Trainer, M. G., Stiles, B., Kirk, R., Hibbitts, C. A., and Malaska, M. J. (2024b). Organic input to Titan’s subsurface ocean through impact cratering. *Astrobiology*, 24(4):378–392. doi:[10.1089/ast.2023.0055](https://doi.org/10.1089/ast.2023.0055).
- Neish, C. D., Somogyi, Á., and Smith, M. A. (2010). Titan’s primordial soup: Formation of amino acids via low-temperature hydrolysis of Tholins. *Astrobiology*, 10(3):337–347. doi:[10.1089/ast.2009.0385](https://doi.org/10.1089/ast.2009.0385).
- Nicolaides, H. (2026). Title of the thesis. Master’s thesis, University of Cambridge.
- Niemann, H. B., Atreya, S. K., Bauer, S. J., et al. (2005). The abundances of constituents of Titan’s atmosphere from the GCMS instrument on the Huygens probe. *Nature*, 438:779–784. doi:[10.1038/nature04122](https://doi.org/10.1038/nature04122).
- Nimmo, F. and Bills, B. G. (2010). Shell thickness variations and the long-wavelength topography of Titan. *Icarus*, 208:896–904. doi:[10.1016/j.icarus.2010.02.020](https://doi.org/10.1016/j.icarus.2010.02.020).
- O’Brien, D. P., Lorenz, R. D., and Lunine, J. I. (2005). Numerical calculations of the longevity of impact oases on Titan. *Icarus*, 173:243–253. doi:[10.1016/j.icarus.2004.08.009](https://doi.org/10.1016/j.icarus.2004.08.009).
- Palmer, M. Y., Cordiner, M. A., Nixon, C. A., Charnley, S. B., Teanby, N. A., Kisiel, Z., Irwin, P. G. J., and Mumma, M. J. (2017). ALMA detection and astrobiological potential of vinyl cyanide on Titan. *Science Advances*, 3(7):e1700022. doi:[10.1126/sciadv.1700022](https://doi.org/10.1126/sciadv.1700022).
- Petricca, F., Vance, S. D., Parisi, M., Buccino, D., Cascioli, G., Castillo-Rogez, J., Downey, B. G., Nimmo, F., Tobie, G., Journaux, B., Lunine, J. I., et al. (2025). Titan’s strong tidal dissipation precludes a subsurface ocean. *Nature*, 648:556–561. doi:[10.1038/s41586-025-09818-x](https://doi.org/10.1038/s41586-025-09818-x).
- Radebaugh, J., Lorenz, R. D., Farr, T., Paillou, P., Savage, C., and Spencer, C. (2011). Linear dunes on Titan and earth: initial remote sensing comparisons. *Geomorphology*, 121:122–132. doi:[10.1016/j.geomorph.2010.04.007](https://doi.org/10.1016/j.geomorph.2010.04.007).
- Rahim, I. S. I. (2026). Title of the thesis. Master’s thesis, University of Cambridge.
- Rahm, M., Lunine, J. I., Usher, D. A., and Shalloway, D. (2016). Polymorphism and electronic structure of polyimine and its potential significance for prebiotic chemistry on Titan. *Proceedings of the National Academy of Sciences*, 113(29):8121–8126. doi:[10.1073/pnas.1606634113](https://doi.org/10.1073/pnas.1606634113).
- Ruiz-Mirazo, K., Briones, C., and de la Escosura, A. (2014). Prebiotic systems chemistry: New perspectives for the origins of life. *Chemical Reviews*, 114(1):285–366. doi:[10.1021/cr2004844](https://doi.org/10.1021/cr2004844).
- Schulze-Makuch, D. and Grinspoon, D. H. (2005). Biologically enhanced energy and carbon cycling on Titan? *Astrobiology*, 5(4):560–567. doi:[10.1089/ast.2005.5.560](https://doi.org/10.1089/ast.2005.5.560).
- Schulze-Makuch, D., Meéndez, A., Fairén, A. G., von Paris, P., Turse, C., Boyer, G., Davila, A. F., Antonio de Sousa António, M. L., et al. (2011). A two-tiered approach to assessing the habitability of exoplanets. *Astrobiology*, 11(10):1041–1052. doi:[10.1089/ast.2010.0592](https://doi.org/10.1089/ast.2010.0592).

- Schwieterman, E. W. and Leung, M. (2024). An overview of exoplanet biosignatures. In Hinkel, N. R., Putirka, K. D., and Xu, S., editors, *Exoplanets: Compositions, Mineralogy, Evolution*, volume 90 of *Reviews in Mineralogy and Geochemistry*, pages 465–514. De Gruyter, Berlin. doi:[10.2138/rmg.2024.90.13](https://doi.org/10.2138/rmg.2024.90.13).
- Seignovert, B., Le Mouélic, S., Brown, R. H., Karkoschka, E., Pasek, V., Sotin, C., and Turtle, E. P. (2019). Titan’s global map combining VIMS and ISS mosaics. doi:[10.22002/D1.1173](https://doi.org/10.22002/D1.1173). URL <https://data.caltech.edu/records/8q9an-yt176>.
- Shannon, C. E. (1948). A mathematical theory of communication. *The Bell System Technical Journal*, 27(3):379–423. doi:[10.1002/j.1538-7305.1948.tb01338.x](https://doi.org/10.1002/j.1538-7305.1948.tb01338.x).
- Shock, E. L. and Holland, M. E. (2007). Quantitative habitability. *Astrobiology*, 7(6):839–851. doi:[10.1089/ast.2007.0098](https://doi.org/10.1089/ast.2007.0098).
- Snyder, J. P. (1987). *Map Projections—A Working Manual*. Number 1395 in U.S. Geological Survey Professional Paper. U.S. Government Printing Office, Washington, D.C. doi:[10.3133/pp1395](https://doi.org/10.3133/pp1395).
- Solomonidou, A. et al. (2018). The spectral nature of Titan’s major geomorphological units: constraints on surface composition. *Journal of Geophysical Research: Planets*, 123:489–507. doi:[10.1002/2017JE005477](https://doi.org/10.1002/2017JE005477).
- Spiegel, D. S. and Turner, E. L. (2012). Bayesian analysis of the astrobiological implications of life’s early emergence on Earth. *Proceedings of the National Academy of Sciences*, 109(2):395–400. doi:[10.1073/pnas.1111694108](https://doi.org/10.1073/pnas.1111694108).
- Stevenson, J., Lunine, J., and Clancy, P. (2015). Membrane alternatives in worlds without oxygen: creation of an azotosome. *Science Advances*, 1(1):e1400067. doi:[10.1126/sciadv.1400067](https://doi.org/10.1126/sciadv.1400067).
- Stofan, E. R., Elachi, C., Lunine, J. I., et al. (2007). The lakes of Titan. *Nature*, 445:61–64. doi:[10.1038/nature05608](https://doi.org/10.1038/nature05608).
- Strobel, D. F. (2010). Molecular hydrogen in Titan’s atmosphere: Implications of the Cassini INMS measurements. *Icarus*, 208(2):878–886. doi:[10.1016/j.icarus.2010.02.009](https://doi.org/10.1016/j.icarus.2010.02.009).
- Tobie, G., Grasset, O., Lunine, J. I., Mocquet, A., and Sotin, C. (2005). Titan’s internal structure inferred from a coupled thermal-orbital model. *Icarus*, 175:496–502. doi:[10.1016/j.icarus.2004.12.007](https://doi.org/10.1016/j.icarus.2004.12.007).
- Tobie, G., Lunine, J. I., and Sotin, C. (2006). Episodic outgassing as the origin of atmospheric methane on Titan. *Nature*, 440:61–64. doi:[10.1038/nature04497](https://doi.org/10.1038/nature04497).
- Tosi, N., Kalousová, K., Spohn, T., et al. (2022). Science goals and new mission concepts for future exploration of Titan’s atmosphere, geology and habitability: POSEIDON (Titan Polar Scout/Orbiter and In Situ Lake Lander and Drone Explorer). *Experimental Astronomy*, 54:1117–1177. doi:[10.1007/s10686-021-09815-8](https://doi.org/10.1007/s10686-021-09815-8).

- Turtle, E. P., Perry, J. E., Hayes, A. G., et al. (2011). Rapid and pervasive changes in surface longevity and coverage observed by Cassini VIMS: Titan’s active methane cycle. *Science*, 331:1414–1417. doi:[10.1126/science.1201063](https://doi.org/10.1126/science.1201063).
- Vuitton, V., Yelle, R. V., and McEwan, M. J. (2007). Ion chemistry and N-containing molecules in Titan’s upper atmosphere. *Icarus*, 191:722–742. doi:[10.1016/j.icarus.2007.06.023](https://doi.org/10.1016/j.icarus.2007.06.023).
- Waite, J. H. et al. (2017). Cassini finds molecular hydrogen in the Enceladus plume: evidence for hydrothermal processes. *Science*, 356:155–159. doi:[10.1126/science.aai8703](https://doi.org/10.1126/science.aai8703).
- Walker, S. I., Bains, W., Cronin, L., DasSarma, S., Danielache, S., Domagal-Goldman, S., Kacar, B., Kiang, N. Y., Lenardic, A., Reinhard, C. T., Moore, W., Schwieterman, E. W., Shkolnik, E. L., and Smith, H. B. (2018). Exoplanet biosignatures: Future directions. *Astrobiology*, 18(6):779–824. doi:[10.1089/ast.2017.1738](https://doi.org/10.1089/ast.2017.1738).
- Wood, C. A. et al. (2010). Impact craters on Titan. *Icarus*, 206:334–344. doi:[10.1016/j.icarus.2009.08.021](https://doi.org/10.1016/j.icarus.2009.08.021).
- Yanez, J. et al. (2024). Acetylenotrophy and the habitability of Titan’s surface. *Icarus*, 408:115969. doi:[10.1016/j.icarus.2024.115969](https://doi.org/10.1016/j.icarus.2024.115969).
- Zebker, H. A., Stiles, B., Hensley, S., et al. (2009). Size and shape of Saturn’s moon Titan. *Science*, 324(5929):921–923. doi:[10.1126/science.1168905](https://doi.org/10.1126/science.1168905).

A Physical Parameters for Titan

Parameter	Value	Reference
Mean radius	2575.0 km	IAU 2015
Surface gravity	1.352 m s^{-2}	Zebker et al. (2009)
Surface pressure	$1.467 \pm 0.001 \text{ bar}$	Fulchignoni et al. (2005)
Surface temperature	$93.65 \pm 0.25 \text{ K}$	Fulchignoni et al. (2005)
N ₂ fraction (surface)	$\approx 94.2\%$	Niemann et al. (2005)
CH ₄ fraction (surface)	$5.65 \pm 0.18\%$	Niemann et al. (2005)
Tidal Love number k_2	0.589 ± 0.150	Iess et al. (2012)
Saturn orbital period	29.46 yr	—
Axial obliquity	26.73°	—
Solar flux at Titan	$\approx 14.9 \text{ W m}^{-2}$	at 9.54 AU
Organic deposition rate	$\sim 5 \times 10^{-14} \text{ g cm}^{-2} \text{ s}^{-1}$	Lavvas et al. (2011)
North polar sea volume	$\sim 70,000 \text{ km}^3$	Mastrogiuseppe et al. (2019)
Water–ammonia eutectic	176 K at 32.6 wt% NH ₃	Grasset and Pargamin (2005)
Subsurface ocean depth	$\sim 55\text{--}80 \text{ km}$	Tobie et al. (2005)

Table A.1: Physical constants and observationally-constrained parameters for Titan used throughout this thesis.

B Bayesian Prior Specification

Feature	Symbol	w_i	μ_i	Key ref.
Liquid hydrocarbon	f_1	0.25	0.020	Hayes et al. (2008)
Organic abundance	f_2	0.20	0.700	Malaska et al. (2025)
Chemical energy	f_3	0.20	0.350	Strobel (2010)
Methane cycle	f_4	0.15	0.400	Mitchell and Lora (2016)
Surface–atm. interaction	f_5	0.08	0.350	Miller et al. (2021)
Topographic complexity	f_6	0.06	0.250	Corlies et al. (2017)
Geomorphological diversity	f_7	0.04	0.300	Lopes et al. (2019b)
Subsurface ocean proximity	f_8	0.02	0.030	Iess et al. (2012)
<i>Global prior mean μ_0</i>			0.331	

Table B.1: Prior means μ_i and weights w_i for all eight habitability features at the present (Cassini) epoch. Global prior mean $\mu_0 = \sum_i w_i \mu_i = 0.331$. Concentration parameter $\kappa = 5$ is applied uniformly.

C Bayesian Inference: Background and Derivation

This appendix provides the conceptual background for the Beta-conjugate Bayesian framework used in Chapter 2. Readers familiar with conjugate Bayesian updating may skip directly to the model summary in Section 2.4.

C.1 The Latent Variable

In this framework, habitability H is treated as a *latent variable* — a quantity that cannot be directly observed by any Cassini instrument but whose value is inferred from the combination of observational proxies $\mathbf{f} = (f_1, \dots, f_8)$. A latent variable (from the Latin *latere*: to be hidden) influences observable outcomes but is not itself directly measurable. Here, H is not a binary yes-or-no property but a continuous probability in $[0, 1]$: the value $H = p$ at a given pixel means that pixel has probability p of satisfying the relevant habitability conditions given all available evidence. Each feature f_i provides partial, indirect evidence about the value of H .

C.2 Why a Beta Distribution?

The Beta distribution $\text{Beta}(\alpha, \beta)$ is defined on $[0, 1]$ and characterised by two positive shape parameters α and β . Its probability density function is expressed in terms of the gamma function (Eq. 2.3 in Chapter 2):

$$p(h) = \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} h^{\alpha-1} (1-h)^{\beta-1}, \quad h \in [0, 1]$$

where the gamma function $\Gamma(\cdot)$ generalises the factorial to non-integer arguments (Eq. 2.4). The distribution is unimodal when both $\alpha > 1$ and $\beta > 1$, with mean $\alpha/(\alpha + \beta)$ and variance that decreases as $\alpha + \beta$ grows. Because it is bounded on $[0, 1]$ and can represent any unimodal belief over a probability, it is the natural prior for the latent habitability probability H .

C.3 Bernoulli Trials and the Binomial Likelihood

A *Bernoulli trial* is an experiment with exactly two possible outcomes — conventionally called “success” and “failure” — each occurring with fixed probabilities p and $1 - p$ respectively (DeGroot and Schervish, 2012, ch. 2). A coin flip is the prototypical example: heads (success, probability p) or tails (failure, probability $1 - p$). The *Binomial distribution* counts the total number of successes k in n independent Bernoulli trials; it is the natural likelihood function for data arising from repeated binary observations.

In the habitability model, no actual coin flips occur. Instead, the prior and data contributions are expressed as *virtual* (or *pseudo*) Bernoulli trials — a mathematical device for encoding strength of belief in the same currency as observational data. The prior contributes $\kappa = 5$ virtual trials, of which $\alpha_0 = 1.655$ are “successes” (prior evidence for habitability) and $\beta_0 = 3.345$ are “failures” (prior evidence against). Each observed feature f_i then contributes λw_i additional virtual trials, of which a fraction f_i are successes and $(1 - f_i)$ are failures. This pseudo-count interpretation is the standard Bayesian device for placing prior beliefs on the same mathematical footing as observed data (Gelman et al., 2013, ch. 2), and is the mechanism by which conjugacy yields a closed-form posterior (Section C.4).

C.4 Conjugacy

A prior distribution is *conjugate* to a likelihood function if the posterior (prior updated by data) has the same mathematical form as the prior, differing only in its parameter values (Gelman et al., 2013, ch. 2). The Beta distribution is conjugate to the Binomial likelihood: if the prior on a success probability p is $\text{Beta}(\alpha_0, \beta_0)$ and we observe k successes in n Bernoulli trials, the posterior is $\text{Beta}(\alpha_0 + k, \beta_0 + (n - k))$. In the habitability model, this update takes the specific form given in Eqs. 2.7–2.8, where the pseudo-counts $k = \lambda \sum_i w_i f_i$ and $n - k = \lambda \sum_i w_i (1 - f_i)$ are derived from the weighted feature values. The resulting posterior PDF (Eq. 2.9) has the same gamma-function structure as the prior PDF (Eq. 2.3), with only the shape parameters changed — no new functional form is introduced. This analytical tractability makes the update feasible at 6.49 million pixels simultaneously with no numerical integration required.

C.5 The Beta-Binomial Model Applied to Habitability

Each feature $f_i \in [0, 1]$ is treated as the fractional outcome of λw_i virtual Bernoulli trials: fraction f_i of those trials are “successes” (evidence for habitability) and $(1 - f_i)$ are “failures” (evidence against). High feature values contribute more positive pseudo-observations and shift the posterior towards higher habitability; low values do the opposite. The prior contributes κ virtual observations and the data from all eight features contribute a further $\lambda \sum_i w_i = \lambda$ observations.

Figure 2.3 visualises the updating process for three representative sites (Freeman Lacus, Belet dune field, Selk crater), showing how each successive feature narrows and displaces the Beta distribution from the prior to the final posterior. The figure also shows the 95% HDI for eight key present-epoch sites.

Figure 2.2 illustrates this update graphically for four representative sites. The left panel shows the prior $\text{Beta}(\alpha_0, \beta_0)$, which is weakly informative ($\sigma \approx 0.19$) and centred at $\mu_0 = 0.331$. The right panel shows the posteriors after incorporating the respective Cassini feature values: Freeman Lacus shifts strongly rightward (highest VIMS organic abundance combined with high liquid and methane-cycle scores), while the Belet dune sea shifts only modestly (moderate organic but no liquid) and Selk crater shifts leftward (below-average methane cycle and surface–atmosphere scores despite exceptional geodiversity).

D Geomorphological Organic Abundance Scores

Class	Name	Physical basis	Score
Dunes	Equatorial sand seas	Tholin-dominated; highest VIMS ratio (Lorenz et al., 2006)	0.82
Plains	Smooth lowlands	Organic sediment blanket (Lopes et al., 2019a)	0.68
Basins	Topographic lows	Organic accumulation (Lopes et al., 2020)	0.58
Labyrinth	Dissected terrain	Spectral affinity with plains (Malaska et al., 2020)	0.55
Craters	Impact structures	Mixed water-ice and organic ejecta (Neish et al., 2018)	0.35
Mountains	Ridge terrain	Water-ice-rich crustal material (Solomonidou et al., 2018)	0.25
Lakes	Hydrocarbon seas	Liquid hydrocarbons; low solid organic content (Brown et al., 2008)	0.05

Table D.1: Organic abundance proxy scores used for the geomorphological blended VIMS mode of Feature 2. Grounded in published VIMS spectral characterisation of each terrain class.

E Declared Assumptions and Known Limitations

- A1. Conditional feature independence.** The eight features are treated as conditionally independent given habitability H . In practice, organic abundance and chemical energy are positively correlated, as are liquid hydrocarbon and surface–atmosphere interaction. The independence assumption is conservative: correlated positive features slightly overstate the posterior, but the framework remains physically motivated and the data-derived importances are consistent with the prior weights (Section 3.4.3).
- A2. Static geomorphological template.** The Lopes et al. (2019) map is used for all five temporal configurations. The Late Heavy Bombardment surface had different dune distributions and crater density, and the future surface will be submerged. The present template is used as the best available proxy, with amplitude modifications applied through epoch-specific prior means and scale functions.
- A3. Blended VIMS organic abundance mode.** The organic abundance feature uses a blended approach: VIMS spectral band ratios where coverage exists ($\sim 50\%$ of the globe), with a global offset correction to align VIMS values to terrain-class scores at the coverage boundary; terrain-class scores are used directly in the remaining hemisphere. This preserves per-pixel spectral variation at polar lake sites while maintaining seamless global coverage.
- A4. Past-epoch Cassini feature contamination.** The past-epoch posterior was computed using present-era Cassini SAR maps, which encode the present polar lake morphology even at $t = -3.5$ Gya. The north polar Bayesian score in the past anchor is inflated above the physically expected prior value. This is mitigated in the multi-epoch reconstruction by using the analytical Bayesian formula directly for pre-lake-formation frames (Section 2.6) and is documented in Section 4.4.
- A5. Multi-epoch temporal scaling accuracy.** The multi-epoch reconstruction applies scalar scale functions $s_i(t)$ to the present-epoch feature maps and re-evaluates the Bayesian posterior formula at each epoch (Section 2.6.1). This is physically motivated but is not equivalent to independent epoch-specific feature extraction. The scale functions are calibrated to reproduce known physical transitions (lake formation, solar warming, eutectic crossing) but have no observational constraint at intermediate epochs. Estimated uncertainty: $\Delta P(H) \lesssim \pm 0.05$ at non-anchor epochs.
- A6. SAR coverage gap.** Approximately 33% of Titan’s surface lacks SAR coverage. Features 1, 3, and 8 are degraded in these regions, where geomorphological and topographic proxies substitute.

F Acronyms, Notation, and Data Source Details

F.1 Acronyms and Nomenclature

The following acronyms are used throughout this chapter and the remainder of the thesis. They are defined here for convenience and also at the point of first use.

SAR	Synthetic Aperture Radar. An active microwave imaging technique in which the motion of the spacecraft is used to synthesise a very long antenna aperture, yielding fine-resolution imagery of surface texture and composition.
VIMS	Visible and Infrared Mapping Spectrometer. The <i>Cassini</i> instrument that imaged Titan’s surface through six near-infrared atmospheric transparency windows, providing surface composition data.
CIRS	Composite Infrared Spectrometer. The <i>Cassini</i> thermal emission instrument used here to constrain the surface temperature field.
GTDR	Global Titan Digital elevation model, Radar — the raw Cassini RADAR altimetry dataset archived in the Planetary Data System.
GTDE	Global Titan Digital Elevation model — the globally interpolated topographic product of Corlies et al. (2017) derived from the GTDR and supplementary data sources.
DEM	Digital Elevation Model. A gridded representation of surface topography.
PDS	Planetary Data System. The NASA archive for planetary mission data.
HDI	Highest Density Interval. A Bayesian credible interval containing the specified probability mass within the narrowest possible range.

F.2 Notation

Two quantities appear throughout this chapter and are defined here to avoid repetition. Each habitability-proxy feature f_i ($i = 1, \dots, 8$) is assigned:

- A **weight** $w_i \in [0, 1]$, where $\sum_i w_i = 1$. The weight encodes the relative astrobiological importance of feature i as a habitability indicator. A feature with a higher weight exerts a larger influence on the posterior probability of habitability, all else being equal. Weights are assigned from physical reasoning and validated against data-derived importances (Section 3.4.3).

- A **prior mean** $\mu_i \in [0, 1]$. This is the expected spatially-averaged value of feature i before any observational data are considered — a statement of our prior knowledge about how common or widespread the habitability-relevant condition is on Titan. For example, liquid hydrocarbon covers $\sim 2\%$ of Titan’s surface (Hayes et al., 2008), so its prior mean is $\mu_1 = 0.02$. The weighted sum of all prior means defines the global prior mean habitability: $\mu_0 = \sum_i w_i \mu_i = 0.331$. Note that the value $\mu_0 = 0.331$ does not appear in the prior literature as a global habitability estimate for Titan; it emerges from the component prior means in Table B.1 and should be interpreted as encoding the observation that Titan’s surface is globally organic-rich whilst simultaneously recognising that liquid solvent is present at only $\sim 2\%$ of the surface. The resulting moderate prior is intentionally revisionable: the weak concentration parameter $\kappa = 5$ ensures that the data, not the prior, determine the posterior at any pixel where observational coverage is good. Full values are tabulated in Appendix B.

This notation first appears in a practical context in Section 2.3, where the eight features are described individually. Their role in the Bayesian update is explained in Section 2.4.

F.3 Observational Data Source Details

F.3.1 Cassini RADAR Synthetic Aperture Radar Mosaic

The Cassini RADAR instrument operated in SAR mode during Titan flybys, imaging the surface at ~ 350 m per pixel over the 127 targeted flybys from 2004 to 2017 (Janssen et al., 2016). SAR measures the normalised radar cross-section σ^0 (typically expressed in dB), which is sensitive to surface roughness, dielectric properties, and the presence of smooth or conducting materials. Low backscatter ($\sigma^0 < -20$ dB) is diagnostic of smooth surfaces (consistent with liquid hydrocarbon lakes) or dielectrically absorbing material (consistent with tholin-rich dune sand). High backscatter correlates with rough, water-ice-rich material such as mountain terrain, crater ejecta, and some cryovolcanic surfaces. The SAR mosaic covers $\sim 67\%$ of Titan’s surface and constitutes the primary data source for the liquid hydrocarbon, chemical energy, and subsurface ocean features.

F.3.2 VIMS Near-Infrared Spectral Mosaic

The Cassini Visible and Infrared Mapping Spectrometer (VIMS) observed Titan’s surface through atmospheric transparency windows at 0.94, 1.08, 1.27, 1.59, 2.03, and 2.79 μm . The band ratio 1.59/1.27 μm is an established proxy for tholin (complex organic) abundance at the surface: the 1.59 μm window probes deeper into the tholin-bearing surface whilst the 1.27 μm window provides atmospheric normalisation (Barnes et al., 2007). The Seignovert et al. (2019) combined VIMS+ISS mosaic (Seignovert et al., 2019) provides this band ratio map at 2–4 km per pixel coverage, but only for the 0–180°W hemisphere ($\sim 50\%$ of the globe). The significance of this coverage gap for the organic abundance feature is discussed in Section G.2.

F.3.3 CIRS Thermal Emission Observations

The Composite Infrared Spectrometer (CIRS) measured Titan’s surface brightness temperature at far-infrared wavelengths throughout the Cassini mission. [Jennings et al. \(2019\)](#) derived an analytical model of the latitude-dependent surface temperature from 13 years of CIRS observations, producing the formula:

$$T(\phi, Y) = (93.53 - 0.095 Y) \cos\left[(\phi + 0.85 - 3.2 Y)(0.0029 - 0.00006 Y)\right] \quad (\text{F.1})$$

where ϕ is latitude in degrees and $Y = t_{\text{CE}} - 2009.61$ is years from the northern spring equinox (22 May 2009). This formula reproduces the observed zonal surface temperatures to within $\pm 0.4 \text{ K}$ (rms) across the full Cassini epoch and provides 100% global coverage, grounded in CIRS focal-plane-1 limb and nadir observations across the full 2004–2017 mission. The model captures only latitudinal and seasonal variation; no longitudinal inhomogeneity is represented (consistent with the absence of significant longitudinal temperature gradients in the Cassini dataset). For temporal epochs other than the present, the mean annual temperature field is held stationary: this approximation is justified because the dominant habitability driver at non-present epochs is the liquid-hydrocarbon scale function $s_1(t)$ rather than the temperature gradient, and the methane-cycle feature contributes only $w_4 = 0.15$ to the posterior. It is used to synthesise the global temperature field and its meridional gradient, which enters the methane cycle feature (Section G.4).

F.3.4 GTDE/GTDR Topographic Models

The Global Titan Digital Elevation model (GTDE), produced by [Corlies et al. \(2017\)](#), integrates Cassini RADAR altimetry, SARTopo stereophotogrammetry, and radial basis function interpolation to produce a globally-interpolated topographic model at $\sim 22 \text{ km}$ per pixel (2 pixels per degree). Coverage is approximately 90% of the globe, with the south polar region and some high-latitude northern areas relying most heavily on interpolation. The topographic model contributes to the chemical energy feature (identifying topographic lows as preferential organic accumulation sites), the topographic complexity feature (through local roughness estimation), and the surface–atmosphere interaction feature (through slope calculation).

F.3.5 Lopes (2019) Global Geomorphological Map

The global geomorphological map of [Lopes et al. \(2019b\)](#) classifies the Cassini SAR-imaged surface into seven terrain units (plains, dunes, mountains, basins, labyrinth terrain, craters, and lakes) using SAR texture, backscatter, morphology, and VIMS spectral correlation. This map provides the terrain classification that enters the organic abundance feature (as a globally seamless gap-fill) and the geomorphological diversity feature. The shapefiles are provided in east-positive geographic coordinates and require conversion to the west-positive convention of the analysis grid.

F.3.6 Miller (2021) Global Fluvial Channel Map

Miller et al. (2021) identified and mapped Titan’s global fluvial channel network from Cassini SAR data, producing a vector dataset of channel locations and widths. This network constitutes the primary mechanism for organic material transport from equatorial source regions to the polar basins, and channel density is used as one component of the surface–atmosphere interaction feature.

F.3.7 Tidal Love Number and Gravity Data

The degree-2 tidal Love number $k_2 = 0.589 \pm 0.150$ from Iess et al. (2012) provides a global constraint on interior structure, confirming the presence of a subsurface liquid layer. This scalar quantity enters as a global floor value in the subsurface ocean proximity feature.

F.3.8 Hedgepeth (2020) Global Crater Catalogue

The global catalogue of confirmed and probable impact structures on Titan by Hedgepeth et al. (2020) records crater diameters and positions for all resolved impact features in the Cassini SAR dataset. This catalogue is used exclusively in the *Past* epoch to generate the impact-melt proximity feature.

F.4 Projection and Reprojection

Three coordinate reference systems are used throughout the pipeline; this section documents each and the transformations between them.

F.4.1 Canonical Equirectangular Grid

All analysis is performed on a simple cylindrical (equirectangular) grid with west-positive longitude (0–360°W) and latitude $\pm 90^\circ$. The grid dimensions are 1802×3603 pixels at 4490 m/px. This projection preserves longitude–latitude coordinates at the cost of area distortion at the poles; it is the native CRS of the Cassini RADAR and GTDR raster products distributed by USGS Astropedia.

F.4.2 GeoTIFF Reprojection via Rasterio

Raw Cassini data products are delivered as GeoTIFFs in a projected coordinate system (coordinates in metres, SimpleCylindrical CRS). These are reprojected onto the canonical grid using `rasterio.warp.reproject` with bilinear resampling. The reprojection handles the coordinate transformation from the source CRS embedded in each GeoTIFF to the canonical grid’s CRS, nodata masking (source nodata values, typically 0.0, are replaced with NaN before reprojection), and resampling from the source resolution to the canonical 4490 m/px grid. This is the only stage at which a GIS reprojection library is used; all subsequent analysis operates in the canonical equirectangular coordinate system.

F.4.3 Polar Stereographic Insets

The polar inset maps in the animation frames and temporal reconstruction figures use an oblique polar stereographic projection following [Snyder \(1987\)](#). The projection maps disc coordinates (x, y) with $r = \sqrt{x^2 + y^2} \in [0, 1]$ to geographic coordinates via:

$$\phi = 90 - 2 \arctan(r \cdot k_{\text{PS}}), \quad \lambda = \arctan 2(x, y) \quad (\text{F.2})$$

where ϕ is latitude, λ is west longitude, and $k_{\text{PS}} = \tan((90 - \phi_{\text{edge}})/2)$ is a pre-computed scale factor that maps the disc edge ($r = 1$) to the desired polar cap boundary ϕ_{edge} (55° by default), $\arctan$ is the single-argument inverse tangent and $\arctan 2$ is the two-argument (quadrant-aware) inverse tangent. For the south polar disc, ϕ is negated. Pixel values are obtained by bilinear interpolation from the equirectangular grid using the inverse-projected latitude and longitude at each disc pixel. This hand-coded implementation avoids introducing a dependency on cartographic projection libraries for what is purely a visualisation transformation.

G Feature Engineering: Scientific Basis and Implementation

The eight habitability-proxy features are extracted from the Cassini data sources described in Appendix F. Each subsection gives the scientific basis, prior mean justification, implementation procedure, and missing-data backfill strategy. Features derived from SAR backscatter (f_1, f_3, f_5, f_7, f_8) are affected by the $\sim 35\%$ areal coverage of the Cassini RADAR SAR imaging mode; in unmapped regions, the pipeline assigns physically motivated floor values rather than interpolated estimates. No published method exists for interpolating SAR backscatter intensity on Titan: unlike topography, which varies smoothly and can be splined from sparse profiles (Corlies et al., 2017; Lorenz et al., 2013), backscatter depends on local roughness, composition, and viewing geometry and lacks the spatial autocorrelation needed for meaningful gap-filling.

Eight habitability-proxy features are extracted from the above data sources. Figure G.1 shows all eight feature maps at the present epoch, illustrating their distinct spatial patterns. Figure G.2 shows how the relative activity of each feature is modulated across geological time.

All features are normalised to the range $[0, 1]$ using a robust percentile-clipping approach:

$$\hat{f}_i = \text{clip}\left(\frac{f_i - f_{i,p2}}{f_{i,p98} - f_{i,p2}}, 0, 1\right) \quad (\text{G.1})$$

where $f_{i,p2}$ and $f_{i,p98}$ are the 2nd and 98th percentiles of finite values in the global layer, mapping the central 96% of the observed dynamic range onto $[0, 1]$ whilst preventing extreme outliers from compressing the scale. Where a feature is unavailable at a particular pixel (due to data gaps), the value is set to NaN and handled through the Bayesian missing-data imputation procedure described in Section 2.4.3.




Figure G.1: All eight habitability-proxy feature maps at the present (Cassini) epoch, each normalised to $[0, 1]$. The panels illustrate the distinct spatial patterns contributed by each feature: the polar concentration of liquid hydrocarbon (f_1), the near-global organic abundance (f_2), the topographically-modulated chemical energy proxy (f_3), and the latitude-dependent methane cycle signal (f_4), amongst others. The combined posterior is derived from all eight simultaneously via the Bayesian update described in Section 2.4. Prior weight w_i and prior mean μ_i are shown for each panel. Blank (white) regions indicate pixels where the feature data is unavailable due to Cassini coverage gaps (primarily in SAR-dependent features f_1, f_3, f_5, f_6, f_8); these pixels are imputed with the feature prior mean μ_i during the Bayesian update (Section 2.4.3), contributing no evidence rather than incorrect evidence. The vertical discontinuity visible in f_3 at 180°W is an assembly seam in the USGS SAR mosaic (c1on180 product), where flyby swaths from opposite hemispheres were stitched at the mosaic’s central meridian; this artefact propagates through the SAR backscatter component of the acetylene energy proxy but is attenuated in the posterior by the 20% weight of f_3 .




Figure G.2: Scale factor (relative activity) for each feature as a function of geological time from -4.0 Gya to $+7.0$ Gya. Colour intensity encodes the scale applied to each feature’s prior mean at that epoch: bright yellow indicates full present-epoch amplitude; dark blue indicates near-zero activity. Key events (LHB peak, lake-formation onset, present epoch, solar warming ramp, eutectic crossing, ocean peak) are marked with vertical dashed lines. The impact-melt proxy and cryovolcanic flux features are active only in pre-lake epochs and are shown in the lower two rows.

G.1 Feature 1: Liquid Hydrocarbon Availability

($w_1 = 0.25, \mu_1 = 0.02$)

Scientific basis. Liquid methane and ethane are the proposed non-aqueous biochemical solvents for any life operating on Titan’s present-day surface (McKay, 2016; McKay and Smith, 2005). At the ambient surface temperature of ~ 94 K, water is structurally as rigid as rock and cannot support molecular diffusion processes fundamental to metabolism. Liquid methane and ethane are, however, stable at 94 K and 1.47 bar, as confirmed by the permanent north polar seas. Benner et al. (2004) proposed that non-polar solvents can in principle support biochemistry through different but self-consistent molecular architectures, and Stevenson et al. (2015) demonstrated through molecular dynamics simulations that acrylonitrile molecules present in Titan’s atmosphere can form stable vesicle membranes in liquid methane. The presence of persistent liquid hydrocarbon is therefore treated as the single most important habitability prerequisite for surface life at the present epoch, and accordingly receives the highest prior weight ($w_1 = 0.25$).

Prior mean. The north polar seas cover approximately 1.5–2% of Titan’s surface (Hayes et al., 2008), giving a prior mean of $\mu_1 = 0.02$. Despite the high weight, the low prior mean means that most pixels contribute negligibly through this feature; its habitability impact is strongly localised to the polar regions.

Implementation. The primary source is the Lopes et al. (2019) lake polygon class (terrain class 7), rasterised to the canonical grid as a binary lake mask. A secondary contribution comes from SAR backscatter dark anomalies not otherwise attributed to organic material. Confirmed empty lake basins from the Birch et al. (2017) analysis are explicitly assigned $f_1 = 0$ to reflect the absence of present-day liquid even within geomorphologically defined basins. In the absence of SAR coverage, ISS 938 nm albedo imagery provides a proxy, with smooth, low-albedo regions consistent with liquid surfaces elevated towards higher feature values.

Backfill. In regions lacking both SAR and ISS coverage, the feature is assigned NaN and imputed at the prior mean during inference.

G.2 Feature 2: Organic Abundance ($w_2 = 0.20, \mu_2 = 0.70$)

Scientific basis. Tholins — the complex refractory macromolecular organic compounds produced continuously by ultraviolet photolysis of the N_2/CH_4 mixture in the upper atmosphere — are the dominant surface material across most of Titan. Their astrobiological significance is multifaceted. Laboratory hydrolysis of tholins in liquid water at biologically relevant temperatures (273–373 K) produces amino acids, nucleobase analogues, and sugar-like polyols at yields of $\sim 0.1\%$ by mass (Neish et al., 2018), directly confirming that tholins can serve as a chemical feedstock for prebiotic synthesis when liquid water is available —

either from impact melts at past epochs or from the red-giant ocean in the far future. Meyer-Dombard et al. (2025) characterise the tholin layer as the primary chemical substrate for any surface habitability scenario on Titan. The feature weight ($w_2 = 0.20$) is second only to liquid hydrocarbon, reflecting this central role.

Prior mean. The revised prior mean of $\mu_2 = 0.70$ (updated from 0.60 in earlier pipeline versions) reflects updated spectral characterisation from Malaska et al. (2025) and Cable et al. (2012), indicating that approximately 82% of the SAR-classified surface by area falls into organic-rich terrain classes (dune: 17%, plains: 65%). The organic abundance feature consequently makes the largest single contribution to the global prior mean ($w_2\mu_2 = 0.140$).

Implementation and the blended VIMS mosaic. The primary spectral data source is the VIMS+ISS 1.59/1.27 μm band-ratio mosaic from Seignovert et al. (2019), but this covers only the 0–180°W hemisphere ($\sim 50\%$ of the globe). Extensive testing of direct blending between the VIMS band ratios and geomorphological class scores in the 180–360°W hemisphere produced a visible seam artefact at the 180°W boundary: the absolute calibration of the VIMS spectral ratio differs systematically from the terrain-class score at the coverage boundary, producing a step change of ~ 0.13 in median organic abundance clearly visible in rendered maps. The analysis resolves this with a global offset correction: the median VIMS value in the overlap region between VIMS and terrain-class coverage is aligned to the corresponding median terrain-class score, eliminating the boundary discontinuity while preserving per-pixel spectral variation within the VIMS-covered hemisphere. In the remaining hemisphere, terrain-class scores from Table D.1 (Appendix D) are used directly. This blended approach retains the spectral resolution of the VIMS data (standard deviation 0.27) at polar lake sites where per-pixel organic variation is scientifically significant, while maintaining seamless global coverage.¹

Backfill. The Lopes (2019) terrain classification covers $\sim 90\%$ of the SAR-imaged surface ($\sim 67\%$ of the globe). In regions outside SAR coverage, terrain class is estimated from ISS albedo units with reduced specificity.

G.3 Feature 3: Chemical Energy Availability ($w_3 = 0.20, \mu_3 = 0.35$)

Scientific basis. Life requires a sustained source of chemical free energy to drive metabolic reactions against thermodynamic equilibrium. On Titan’s surface, the primary candidate is acetylenotrophy ($\text{C}_2\text{H}_2 + 3\text{H}_2 \longrightarrow 2\text{CH}_4$, $\Delta G = -334\text{ kJ mol}^{-1}$; where ΔG is the Gibbs

¹The Seignovert VIMS+ISS mosaic uses IAU east-positive longitude (Seignovert et al. 2019; confirmed by the 2019 LPSC abstract referencing landmarks at east-positive coordinates). The pipeline canonical grid uses west-positive longitude. Conversion requires a horizontal flip of the pixel array (not a 180° roll), because the two conventions differ in column ordering direction, not origin. Alignment was validated by three tests: (i) Kraken Mare (310°W) shows VIMS spectral texture (local $\sigma > 0.02$); (ii) no step change at the 180°W hemisphere boundary ($|\Delta P(H)| < 0.02$ at all latitude bands); and (iii) the rendered global maps show smooth spatial gradients consistent with the Cassini-era geomorphological observations.

Free Energy change of the reaction, a negative value indicating spontaneous energy release) available from acetylene–hydrogen combination (McKay and Smith, 2005). Yanez et al. (2024) recalculated the available energy under updated thermochemical conditions at Titan’s surface as 69–78 kJ mol^{−1} C, well above the minimum for terrestrial methanogens. Acetylene is produced continuously by UV photolysis of methane and sediments to the surface where it has been directly detected by GCMS (Niemann et al., 2005). The chemical energy feature is therefore a proxy for the co-availability of acetylene and hydrogen at the surface.

Implementation. Since the acetylene surface abundance cannot be directly mapped by any Cassini instrument, it is approximated through two components. SAR backscatter (weight 70% where coverage exists): SAR-dark surfaces characteristic of organic-rich terrain represent regions where acetylene is most likely to accumulate without competing abiotic reactions consuming it rapidly. Topographic inversion (weight 30%): acetylene and other atmospheric condensates preferentially accumulate in topographic lows through both gravitational settling and fluvial transport; the normalised negative of the GTDE elevation ($-z$) assigns high values to topographic lows. In SAR-absent regions, the feature reverts to the pure topographic inversion component.

G.4 Feature 4: Methane Cycle Intensity

($w_4 = 0.15$, $\mu_4 = 0.40$)

Scientific basis. The methane cycle is the dominant agent of chemical transport on Titan. It delivers organic material from equatorial photochemical production regions to the polar lake margins, maintains chemical gradients at shorelines, and provides seasonal wet–dry cycling at lake edges that some prebiotic chemistry scenarios identify as important for concentrating organic reactants and driving condensation reactions (Benner et al., 2004; Mitchell and Lora, 2016). Regions experiencing the most active methane cycling receive fresh organic inputs most regularly and maintain the most dynamic chemical environments.

Implementation. The methane cycle intensity is a weighted composite of three independently computed components:

1. **VIMS observation density** (50%): the number of VIMS spectral footprints per grid cell, smoothed with a 5-pixel Gaussian kernel. Cassini’s VIMS observation targeting was scientifically motivated, prioritising regions of active methane-cycle interest; higher coverage density therefore correlates with active methane cycling zones (Le Mouélic et al., 2019).
2. **Meridional temperature gradient** $|\partial T / \partial \phi|$ (25%): derived from the Jennings (2019) model (Eq. F.1), evaluated at $Y = +1.39$ (year 2011.0, the approximate mission midpoint). Regions where the surface temperature gradient is steepest experience the strongest differential evaporation and precipitation, driving the most vigorous methane transport.

3. **Latitude prior** (25%): the function $\exp[-(\phi \pm 45)^2/(25)^2]$, following [Mitchell and Lora \(2016\)](#)’s general circulation modelling result that the mid-latitude regions host the primary methane cycle activity during the Cassini epoch.

This feature has 100% global coverage by construction and requires no backfill.

G.5 Feature 5: Surface–Atmosphere Interaction Intensity ($w_5 = 0.08, \mu_5 = 0.35$)

Scientific basis. The physical interface between surface and atmosphere is where gas-phase photochemical products encounter potential liquid solvents and organic substrates. Lake margins and fluvial channel termini are the primary loci of this exchange. [Hayes \(2016\)](#) describes the complex sedimentological and chemical processes at hydrocarbon lake shorelines, including the evaporation-driven concentration of dissolved organics and the seasonal flooding and drying cycles that expose organic sediments to liquid contact. [Miller et al. \(2021\)](#) documents that the fluvial channel network terminates preferentially at polar lake margins, establishing these as zones of maximum organic delivery and chemical gradient activity.

Implementation. A weighted sum of three spatial components:

$$f_5 = 0.30 \hat{s} + 0.40 \hat{m} + 0.30 \hat{c} \quad (\text{G.2})$$

where $\hat{s}$ is the normalised topographic slope $|\nabla z|$ from the GTDE DEM (high slopes drive physical transport and mixing); $\hat{m}$ is a binary lake-margin mask extended to ~ 45 km from the nearest confirmed lake shore; and $\hat{c}$ is the Gaussian-smoothed ($\sigma = 3$ pixels) channel density from the [Miller et al. \(2021\)](#) dataset, normalised at the 2nd–98th percentile range. The lake-margin component receives the highest within-feature weight (40%) because shorelines concentrate all three relevant processes — liquid contact, organic delivery, and thermal/chemical cycling.

G.6 Feature 6: Topographic Complexity ($w_6 = 0.06, \mu_6 = 0.25$)

Scientific basis. Local topographic roughness at ~ 20 km scale correlates with the dissected terrain types that most frequently expose water-ice-bearing crustal material ([Lopes et al., 2019b](#)). On Earth, topographic heterogeneity at ecologically relevant scales creates micro-environmental diversity — variations in drainage, aspect, moisture, and chemical exposure — that underpins elevated species richness at landscape boundaries. On Titan, the analogous principle is that varied topography produces diverse chemical micro-environments with potentially richer reactive diversity.

Implementation. Local standard deviation of GTDE elevation computed in a 5×5 pixel (~ 22 km) sliding window, normalised to $[0, 1]$ at the 2nd–98th percentile range. South polar gaps in the primary GTDE are filled with the Corlies et al. (2017) globally interpolated DEM prior to roughness computation.

G.7 Feature 7: Geomorphological Diversity ($w_7 = 0.04$, $\mu_7 = 0.30$)

Scientific basis. Astrobiologically relevant chemistry is disproportionately likely to occur at interfaces between geochemically distinct terrain types. The reasoning is thermodynamic: chemical gradients — differences in redox potential, organic content, water-ice availability, or pH — drive reactions that would not proceed in chemically uniform settings (Shock and Holland, 2007). On Titan, the most important geochemical interfaces identified from Cassini data are:

1. The **dune–plains boundary**, where organic-rich tholin-saturated dune sand is in direct lateral contact with the moderate-organic plains through which fluvial channels transport material towards the polar lakes. This interface concentrates the exchange between the primary photochemical organic source (dunes) and the active transport system (plains channels).
2. The **crater ejecta interface**, where water-ice-rich crustal material excavated by impact is juxtaposed with adjacent organic-bearing terrain. Neish et al. (2018) demonstrated experimentally that mixing water ice and tholins under liquid water conditions produces amino acids; the spatial co-occurrence of these two materials at crater ejecta boundaries is therefore a direct proxy for the relevant chemical reactivity.
3. The **lake–plains/dune boundary**, where liquid hydrocarbon meets the organic substrate, enabling dissolution and transport of surface organics into the lake system.

The Shannon entropy of terrain class distribution within a local window quantifies the degree of terrain-type mixing at each pixel. A pixel surrounded by a single terrain class (entropy = 0) sits in a geochemically homogeneous environment. A pixel at the intersection of four terrain classes (entropy approaching $H_{\max}$) sits at a multi-way geochemical interface. This feature receives the lowest weight of the present-epoch set ($w_7 = 0.04$) because it is an indirect and geometrical proxy for chemical gradients rather than a direct measurement.

Implementation. The Shannon entropy of the terrain class distribution in a sliding window of radius 10 pixels (~ 45 km), normalised by the theoretical maximum entropy for the number of classes present:

$$f_7(r, c) = \frac{H(r, c)}{H_{\max}}, \quad H(r, c) = - \sum_{k=1}^{N_c} p_k(r, c) \ln p_k(r, c), \quad H_{\max} = \ln N_c \quad (\text{G.3})$$

where $p_k(r, c)$ is the fraction of window pixels classified as terrain class k , and N_c is the number of distinct classes present in the window. Shannon entropy is the standard information-theoretic measure of diversity (Shannon, 1948) and has been applied to quantify geomorphological and ecological diversity in planetary and terrestrial settings (Lorenz, 2014). The normalisation by $H_{\max} = \ln N_c$ maps the maximum possible entropy (uniform distribution across all present classes) to 1.0, ensuring the feature is comparable across pixels with different numbers of terrain classes.

G.8 Feature 8: Subsurface Ocean Proximity

$(w_8 = 0.02, \mu_8 = 0.03)$

Scientific basis. Titan’s confirmed subsurface water–ammonia ocean (Iess et al., 2012) is potentially habitable by conventional criteria (Affholder et al., 2025), and there is indirect evidence that material from this ocean has reached the surface at specific locations. SAR-bright annular features surrounding several impact craters are most consistently interpreted as ancient impact-melt or cryolava fronts where liquid spread outward from basin centres, potentially establishing chemical contact between the subsurface ocean and the organic surface inventory (Wood et al., 2010). The prior mean is set at $\mu_8 = 0.03$ (revised downward from 0.10 following Neish et al. 2024a), reflecting the genuine uncertainty in the surface expression of the subsurface ocean at the present epoch.

Implementation. The feature is constructed in two parts. First, a global floor value of $p_{\text{floor}} = 0.03$ is applied uniformly to all pixels, reflecting the global presence of the subsurface ocean as confirmed by the tidal Love number k_2 (Iess et al., 2012). This floor captures the fact that the ocean is a global structure and that the possibility of ocean–surface exchange cannot be excluded at any location, even if no surface expression is visible.

Second, localised enhancements above the floor are applied at SAR-bright annular structures identified in the Cassini RADAR mosaic. These annuli are selected by their backscatter ($\sigma^0 > 0$ dB, indicating rough or water-ice-rich material), their plan-view morphology (roughly circular or lobate, distinct from the linear patterns of mountain ridges or the radar-dark texture of dunes), and their association with confirmed or probable impact structures in the Hedgepeth et al. (2020) catalogue. At each such structure, a two-dimensional Gaussian kernel of half-width $\sigma \approx 200$ km centred on the crater creates a smooth spatial enhancement. The 200 km scale is set to approximate the spatial extent of the ancient melt-water or cryolava deposit that the annulus represents: for a 90 km diameter crater (Selk), melt-front runout distances of ~ 200 –500 km are consistent with the low-viscosity water–ammonia melt and the low-gravity environment (O’Brien et al., 2005).

The feature is the normalised sum of the global floor and all Gaussian kernels, clipped to $[0, 1]$. Because the global floor is already low ($\mu_8 = 0.03$) and only a small number of confirmed annuli exist in the SAR mosaic, the spatial variance of f_8 is very low: it is 0.03 everywhere except within ~ 500 km of a confirmed annular structure, where it reaches values of 0.15–0.25.

H The Cassini–Huygens Mission and Dragonfly

The *Cassini–Huygens* mission (2004–2017) provided the observational foundation of this thesis through 127 targeted Titan flybys, the *Huygens* atmospheric probe descent and landing, and a systematic multi-instrument survey of Titan’s surface, atmosphere, and interior. SAR mapping covered $\sim 67\%$ of the surface; VIMS near-infrared spectroscopy $\sim 50\%$; CIRS thermal mapping the full globe at $\sim 10^\circ$ resolution. The mission’s key unresolved gap is the absence of a quantitative, spatially-resolved, multi-epoch habitability analysis integrating all data modalities — which this thesis provides.

NASA’s *Dragonfly* rotorcraft ([Lorenz et al., 2018](#)), due to arrive in 2034, will traverse ~ 175 km from the Belet dune field to Selk crater over a 2.7-year mission. Chapter 3 provides a quantitative evaluation of that site selection relative to globally optimal alternatives.

I Dragonfly Mission: Site Validation and Instrument Mapping

I.1 Mission Science Objectives

NASA’s *Dragonfly* rotorcraft (Lorenz et al., 2018), due at Titan in 2034, traverses ~ 175 km from Belet dunes to Selk crater. Its instruments — DraMS (mass spectrometry), DraGNS (gamma-ray and neutron spectrometry), DragonCam (imaging), and DraGMet (meteorology) — directly measure the features underpinning the habitability scores at its landing sites.

I.2 Reconciling Selk’s Low Score with Mission Selection

RQ4 asks whether the habitability map independently validates the *Dragonfly* landing site at Selk crater. Selk scores $P(H) = 0.243$ — well below the prior mean of 0.331 and ~ 35 percentage points below Freeman Lacus. On the face of it, this appears to *contradict* the site selection. The resolution lies in recognising that this thesis and the *Dragonfly* mission address fundamentally different habitability questions.

This thesis: present-day solvent habitability. The pipeline asks: *where does liquid solvent currently exist that could support metabolic chemistry?* The eight features weight liquid hydrocarbon availability ($w_1 = 0.25$), methane cycle activity ($w_4 = 0.15$), and surface-atmosphere interaction ($w_5 = 0.08$) — all present-day physical processes. By construction, polar lake shores score highest ($P(H) \approx 0.541$; Freeman Lacus), and dry equatorial craters score low. Selk’s low score is therefore the *expected* output of a correctly functioning present-day solvent model.

Dragonfly: prebiotic chemistry habitability. The mission asks: *where has liquid water most plausibly contacted organic material?* Neish et al. (2018) showed that large impact craters are the primary targets: a major impact melts the water-ice crust, producing transient liquid water (10^3 – 10^6 yr; Artemieva and Lunine 2003; O’Brien et al. 2005) that drives Strecker synthesis of amino acids and nucleobase precursors from surface tholins (Neish et al., 2010). *Dragonfly* explicitly targets *evidence of past water-organic contact*, not present-day liquid (Lorenz et al., 2021).

The two frameworks give opposite site rankings because the polar lakes contain non-polar hydrocarbon liquid that cannot drive the ionic chemistry required for Strecker synthesis, while equatorial craters lack present-day liquid but preserve evidence of past aqueous processing. Neish et al. (2018) explicitly concluded that the polar lakes are *not* the best places to search for biosignatures. The pipeline’s low score for Selk is therefore not a contradiction

of the mission rationale — it is a *consequence* of the two frameworks targeting different habitability questions. The independent validation of Selk comes not from its aggregate $P(H)$ but from a specific feature within the pipeline, as shown in the following section.

I.3 Quantitative Validation of Site Selection

The pipeline is a present-day solvent habitability model. By this metric, the north polar lake margins are the most habitable environments on Titan, reaching $P(H) \approx 0.541$ (Freeman Lacus) and 0.51 (Oib Lacus). Selk scores $P(H) = 0.243$, below the prior mean of 0.33, because it lacks present-day surface liquid ($f_1 \approx 0.01$) and its methane-cycle and surface-atmosphere scores are equatorial-low. *This is scientifically correct*: a model that correctly captures present-day solvent habitability should rank polar lake shores above a dry equatorial crater.

The independent validation comes from the geomorphological diversity feature. Selk's $f_7 = 0.659$ is the highest value of any equatorial site and $7.2\times$ the global median of 0.092 — a result derived from Cassini SAR terrain data alone, without reference to mission planning. Geomorphological diversity captures the simultaneous presence of four terrain classes (dune, plains, labyrinth, crater floor) within 45 km, which is the Cassini-era observational proxy for the organic–water-ice mixing interface that defines the prebiotic chemistry target. The pipeline independently identifies Selk as the equatorial site with the highest terrain-class heterogeneity — precisely the physical attribute cited in the *Dragonfly* science rationale (Lorenz et al., 2021; Neish et al., 2018).

The polar lake margins represent a complementary, not competing, science target. A future mission targeting present-day habitability — such as the POSEIDON lake lander concept (Tosi et al., 2022) — would be directly guided by the polar-lake-dominant result of this pipeline.

I.4 Instrument–Feature Correspondence

DraMS atmospheric and surface composition measurements constrain f_3 (acetylene abundance) through depletion relative to photochemical model predictions, and f_2 (organic inventory) via amino acid and HCN polymer detection. DraGNS measures bulk hydrogen abundance, constraining f_8 (subsurface water proximity). The DraGMet EFIELD experiment directly measures the surface electric field, constraining f_4 .

Enantiomeric excess in amino acids (expected $>10\%$ if biogenic vs. racemic for abiotic Strecker products) is the highest-priority biosignature for DraMS. Acetylene depletion below the photochemical steady-state prediction (McKay and Smith, 2005) would constitute strong but non-unique evidence for acetylenotrophy. Carbon isotopic fractionation ($\delta^{13}\text{C}$ depletion of $\sim 30\text{--}50\%$ relative to atmospheric CH_4) would be consistent with biological methane cycling. All three biosignatures have plausible abiotic explanations and must be interpreted jointly.

I.5 Interpreting a Null Result

A null biosignature result at Selk would constrain specifically the “preserved impact-melt Strecker chemistry” pathway at that location and epoch. It would not constrain: (1) acetylenotrophy at the north polar lake margins, which score ~ 30 percentage points higher than Selk in the present epoch and which *Dragonfly* does not visit; (2) life in the subsurface ocean, insensitive to surface sampling; or (3) extinct life from the LHB whose signatures may be buried below accessible depth. The habitability maps provide a rigorous framework for reporting a null result in terms of which specific pathways have been tested and which remain open.

J Comparison with Prior Habitability Frameworks

Several preceding studies have attempted to place planetary habitability on a quantitative footing. The *Earth Similarity Index* (Schulze-Makuch et al., 2011) scores planetary bodies against Earth on temperature, density, escape velocity, and surface chemistry, but produces a single scalar per body and cannot represent spatial or temporal variation. Hendrix et al. (2019) reviewed the habitability potential of icy ocean worlds using qualitative multi-criteria matrices, providing rich scientific context but no spatial resolution or probabilistic framework. Logistic regression and machine-learning habitability classifiers (e.g. Spiegel and Turner 2012) have been applied to exoplanet populations but require labelled training data that does not exist for Titan.

The Bayesian approach adopted in this thesis differs from all of these: it encodes scientific prior knowledge as probability distributions, updates them with observational data via conjugate updating, and returns credible intervals rather than point estimates — making uncertainty explicit rather than implicit. The choice of a Beta-conjugate (Beta-Binomial) model is motivated by the bounded $[0,1]$ domain of habitability probability and by the analytical tractability that permits simultaneous evaluation at all 6.49 million grid pixels without numerical integration or MCMC. A Gaussian process or multivariate logistic regression could in principle capture feature correlations more richly, but would require either strong covariance priors or prohibitively dense labelled sampling to estimate the correlation structure from Cassini data alone. The Beta-conjugate approach is preferred as the most tractable framework that preserves physical interpretability at every pixel.

The closest methodological precedent is Affholder et al. (2021), who applied Bayesian statistical inference to Cassini INMS plume data from Enceladus, comparing the observed methane and hydrogen escape rates against models of abiotic serpentinisation and biotic methanogenesis. Their finding — that the observed methane escape rate is most parsimoniously explained by biological methanogenesis if the prior probability of life is sufficient — represents the most quantitatively rigorous Bayesian biosignature assessment yet published for a solar system body. The present work is directly inspired by this approach, extended from a single-point plume observation to a full-globe, multi-feature, multi-epoch spatial mapping.